




# Latent Structure Models for NLP

**Tsvetomila Mihaylova** Instituto de Telecomunicações  
**Vlad Niculae** Instituto de Telecomunicações

*work with:*

André Martins Instituto de Telecomunicações & IST & Unbabel  
Nikita Nangia NYU

 [deep-spin.github.io/tutorial](https://deep-spin.github.io/tutorial)

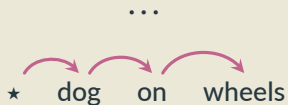
# I. Introduction

# Structured prediction and NLP

- **Structured prediction:** a machine learning framework for predicting structured, constrained, and interdependent outputs
- **NLP** deals with *structured* and *ambiguous* textual data:
  - machine translation
  - speech recognition
  - syntactic parsing
  - semantic parsing
  - information extraction
  - ...

# Examples of structure in NLP

## Dependency parsing



...

# Examples of structure in NLP

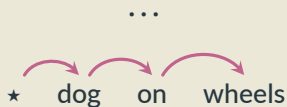
## POS tagging

VERB    PREP    NOUN  
dog    on    wheels


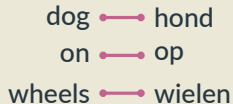
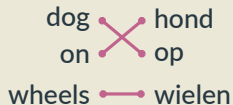
NOUN    PREP    NOUN  
dog    on    wheels

NOUN    DET    NOUN  
dog    on    wheels

## Dependency parsing

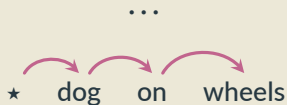


## Word alignments



# Examples of structure in NLP

## Dependency parsing



Exponentially many structures!



...

# **NLP 5 years ago:**

## Structured prediction and pipelines



# NLP 5 years ago:

## Structured prediction and pipelines

- Big pipeline systems, connecting different structured predictors, trained separately
- **Advantages:** fast and simple to train, can rearrange pieces 😊



# NLP 5 years ago:

## Structured prediction and pipelines

- Big pipeline systems, connecting different structured predictors, trained separately
- **Advantages:** fast and simple to train, can rearrange pieces 😊
- **Disadvantage:** linguistic annotations required for each component 😓

# NLP 5 years ago:

## Structured prediction and pipelines

- Big pipeline systems, connecting different structured predictors, trained separately
- **Advantages:** fast and simple to train, can rearrange pieces 😊
- **Disadvantage:** linguistic annotations required for each component 😓
- **Bigger disadvantage:** error propagates through the pipeline 💩

# NLP today:

## End-to-end training



# NLP today:

## End-to-end training

- Forget pipelines—train everything from scratch!
- No more error propagation or linguistic annotations! 🎉

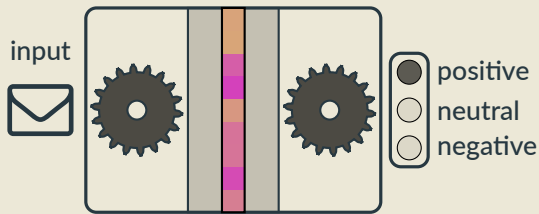
# NLP today:

## End-to-end training

- Forget pipelines—train everything from scratch!
- No more error propagation or linguistic annotations! 🎉
- Treat everything as *latent*! 🙌

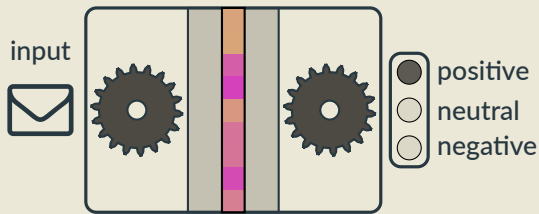
# Representation learning

- Uncover hidden representations useful for the *downstream task*.
- Neural networks are well-suited for this: *deep computation graphs*.



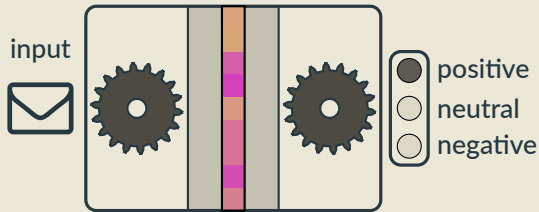
# Representation learning

- Uncover hidden representations useful for the *downstream task*.
- Neural networks are well-suited for this: *deep computation graphs*.
- Neural representations are unstructured, inscrutable.  
Language data has underlying structure!



# Latent structure models

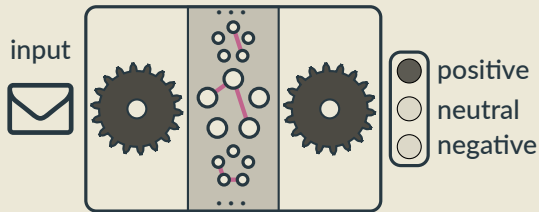
- Seek *structured* hidden representations instead!






# Latent structure models

- Seek *structured* hidden representations instead!




# Latent structure models

- Seek *structured* hidden representations instead!
- They can bring us:
  - More interpretability;




# Latent structure models

- Seek *structured* hidden representations instead!
- They can bring us:
  - More interpretability;
  - Better inductive bias;



# Latent structure models

- Seek *structured* hidden representations instead!
- They can bring us:
  - More interpretability;
  - Better inductive bias;
  - Hopefully: smaller models.



# Latent structure models aren't so new!

- They have a very long history in NLP:
  - IBM Models for SMT (latent word alignments) [Brown et al., 1993]
  - HMMs [Rabiner, 1989]
  - CRFs with hidden variables [Quattoni et al., 2007]
  - Latent PCFGs [Petrov and Klein, 2008, Cohen et al., 2012]
- Trained with EM, spectral learning, method of moments, ...
- Often, very strict assumptions (e.g. strong factorizations)
- Today, neural networks opened up some new possibilities!

# What this tutorial is about:

- Discrete, combinatorial latent structures
- Often the structure is inspired by some linguistic intuition
- We'll cover both:
  - RL methods (structure built incrementally, reward coming from downstream task)
  - ... vs end-to-end differentiable approaches (global optimization, marginalization)
  - stochastic computation graphs
  - ... vs deterministic graphs.
- All plugged in *discriminative* neural models.

# This tutorial is *not* about:

- It's not about continuous latent variables
- It's not about deep generative learning
- We won't cover GANs, VAEs, etc.
- There are (very good) recent tutorials on deep variational models for NLP:
  - “Variational Inference and Deep Generative Models” (Schulz and Aziz, ACL 2018)
  - “Deep Latent-Variable Models for Natural Language” (Kim, Wiseman, Rush, EMNLP 2018)

**Background**



# Unstructured vs structured

- Simplest example of structure: Just a discrete choice among  $N$  categories.
- We call this *unstructured*.
- It will provide an important starting point.

# The challenge of discrete choices

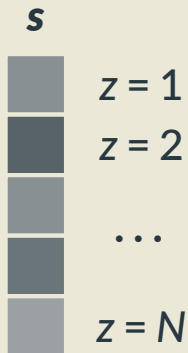
$$z = 1$$

$$z = 2$$

...

$$z = N$$

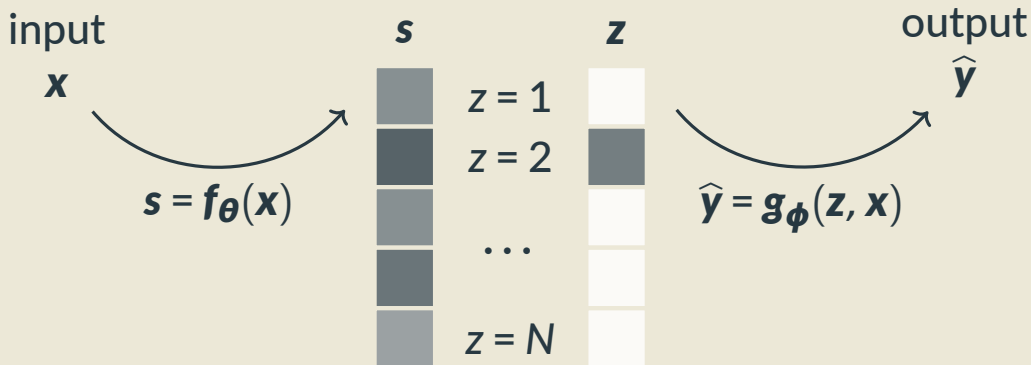
# The challenge of discrete choices



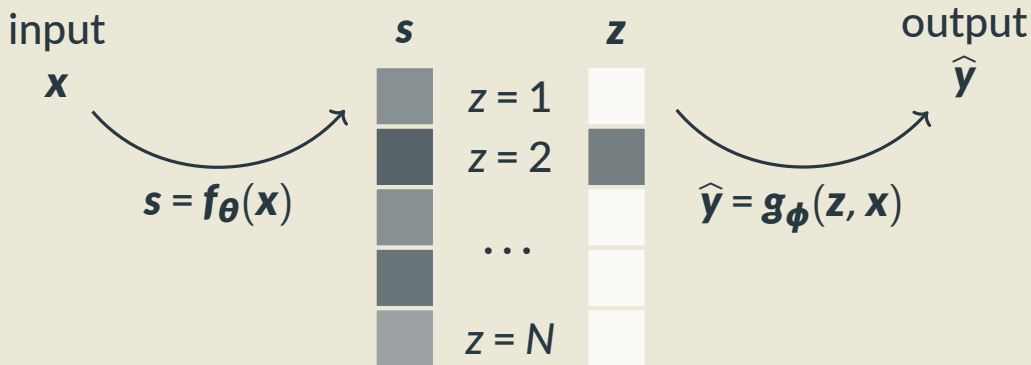
# The challenge of discrete choices



# The challenge of discrete choices

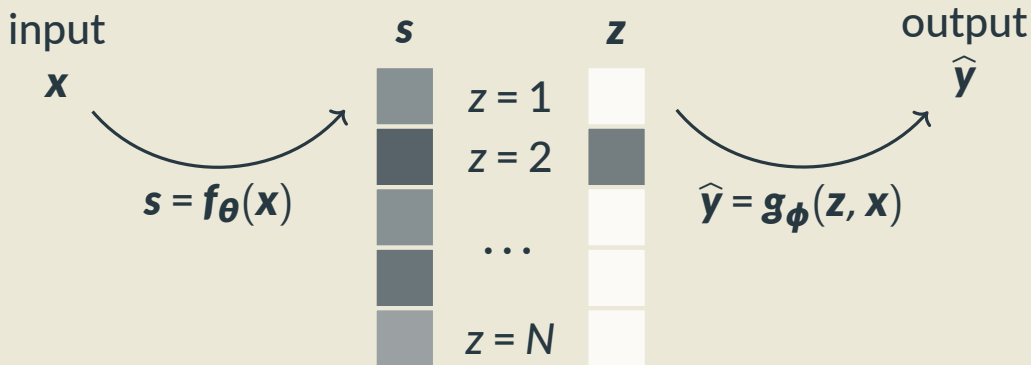


# The challenge of discrete choices



$$\frac{\partial L(\hat{y}, y)}{\partial w} = ?$$

# The challenge of discrete choices

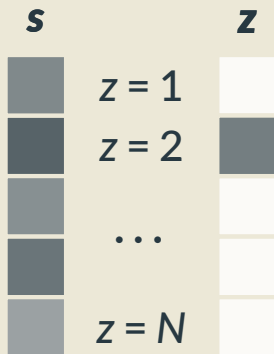


$$\frac{\partial L(\hat{\mathbf{y}}, \mathbf{y})}{\partial \mathbf{w}} = ?$$

or, essentially,

$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = ?$$

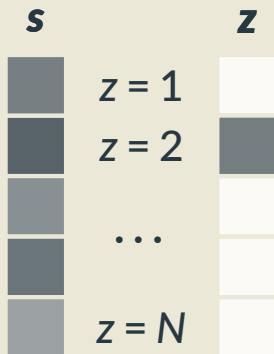
# Discrete mappings are “flat”



$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = ?$$

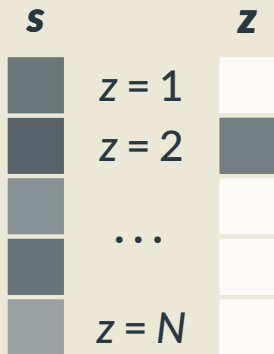


# Discrete mappings are “flat”



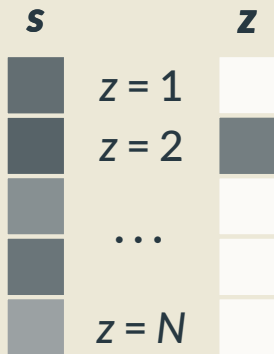
$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = ?$$

# Discrete mappings are “flat”



$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = ?$$

# Discrete mappings are “flat”



$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = ?$$

# Discrete mappings are “flat”



$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = ?$$

# Discrete mappings are “flat”



$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = ?$$

# Discrete mappings are “flat”



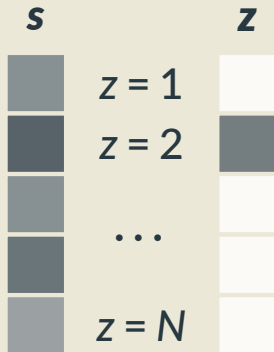
$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = ?$$

# Discrete mappings are “flat”

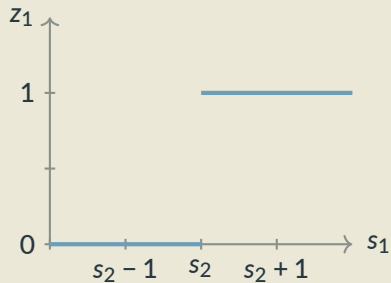


$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = ?$$

# Argmax



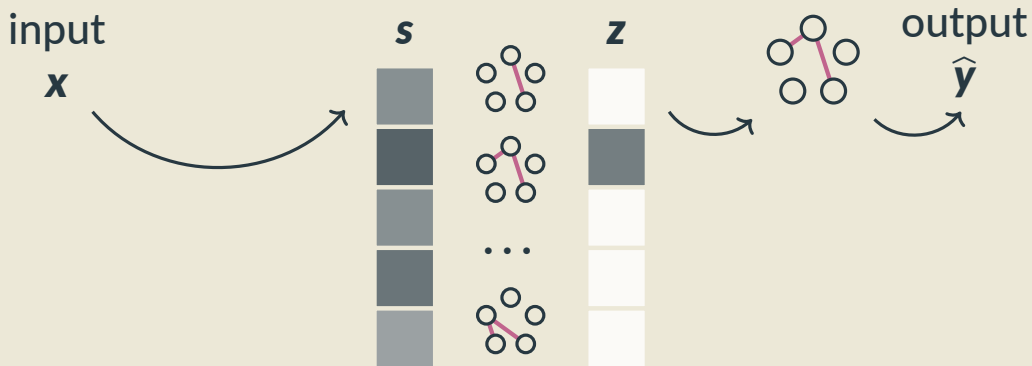
$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = \mathbf{0}$$





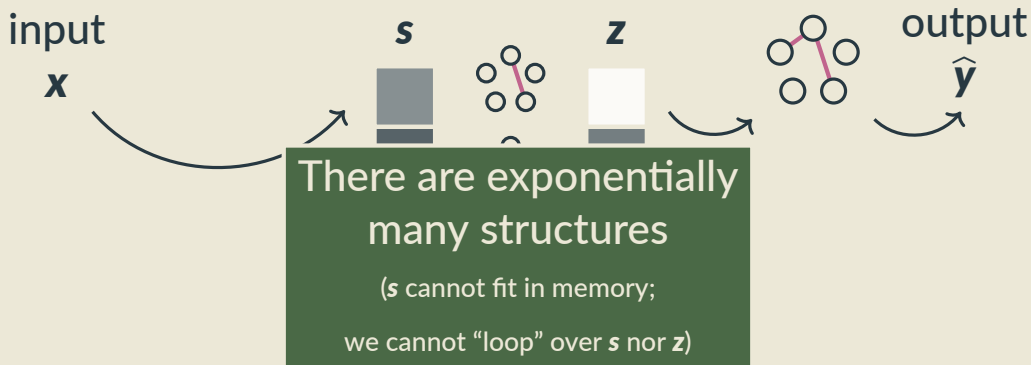
# Computing the most likely structure

is a very high-dimensional argmax

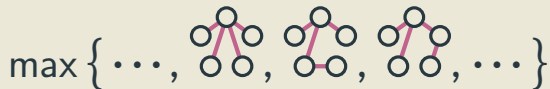
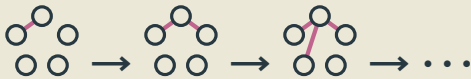


# Computing the most likely structure

is a very high-dimensional argmax



# Dealing with the combinatorial explosion



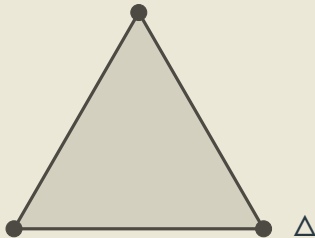
## 1. Incremental structures

- Build structure **greedily**, as sequence of discrete choices (e.g., shift-reduce).
- Scores (partial structure, action) tuples.
- **Advantages:** flexible, rich histories.
- **Disadvantages:** greedy, local decisions are suboptimal, error propagation.

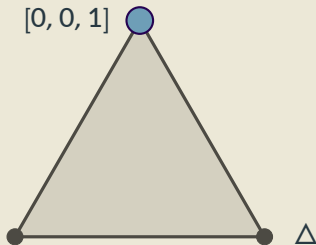
## 2. Factorization into parts

- Optimizes **globally** (e.g. Viterbi, Chu-Liu-Edmonds, Kuhn-Munkres).
- Scores smaller parts.
- **Advantages:** optimal, elegant, can handle hard & global constraints.
- **Disadvantages:** strong assumptions.

# The unstructured case: Probability simplex



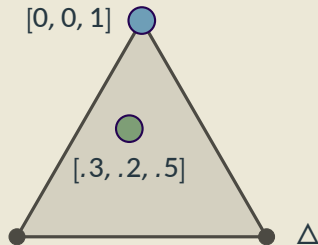
# The unstructured case: Probability simplex



- Each vertex is an *indicator vector*, representing one class:

$$\mathbf{z}_c = [0, \dots, 0, \underbrace{1}_{c^{\text{th}} \text{ position}}, 0, \dots, 0].$$

# The unstructured case: Probability simplex



- Each vertex is an *indicator vector*, representing one class:

$$\mathbf{z}_c = [0, \dots, 0, \underbrace{1}_{c^{\text{th}} \text{ position}}, 0, \dots, 0].$$

- Points inside are *probability vectors*, a convex combination of classes:

$$\mathbf{p} \geq \mathbf{0}, \quad \sum_c p_c = 1.$$

# What's the analogous of $\Delta$ for a structure?

- A structured object  $\mathbf{z}$  can be represented as a *bit vector*.

# What's the analogous of $\Delta$ for a structure?

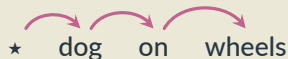
- A structured object  $\mathbf{z}$  can be represented as a *bit vector*.
- Example:
  - a dependency tree can be represented a  $O(L^2)$  vector indexed by arcs
  - each entry is 1 iff the arc belongs to the tree
  - **structural constraints:** not all bit vectors represent valid trees!



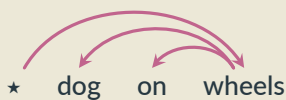
# What's the analogous of $\Delta$ for a structure?

- A structured object  $\mathbf{z}$  can be represented as a *bit vector*.
- Example:
  - a dependency tree can be represented a  $O(L^2)$  vector indexed by arcs
  - each entry is 1 iff the arc belongs to the tree
  - **structural constraints:** not all bit vectors represent valid trees!

$$\mathbf{z}_1 = [1, 0, 0, 0, 1, 0, 0, 0, 1]$$



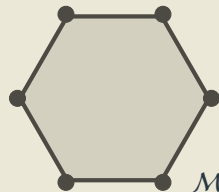
$$\mathbf{z}_2 = [0, 0, 1, 0, 0, 1, 1, 0, 0]$$



$$\mathbf{z}_3 = [1, 0, 0, 0, 1, 0, 0, 1, 0]$$




# The structured case: Marginal polytope



# The structured case: Marginal polytope

- Each vertex corresponds to one such *bit* vector  $\mathbf{z}$




# The structured case: Marginal polytope

- Each vertex corresponds to one such *bit vector*  $\mathbf{z}$
- Points inside correspond to *marginal distributions*: convex combinations of structured objects

$$\boldsymbol{\mu} = \underbrace{p_1 \mathbf{z}_1 + \dots + p_N \mathbf{z}_N}_{\text{exponentially many terms}}, \quad \mathbf{p} \in \Delta.$$

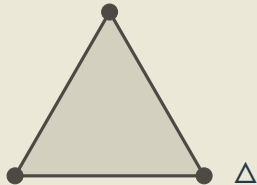
$$\begin{aligned} p_1 &= 0.2, & \mathbf{z}_1 &= [1, 0, 0, 0, 1, 0, 0, 0, 1] \\ p_2 &= 0.7, & \mathbf{z}_2 &= [0, 0, 1, 0, 0, 1, 1, 0, 0] \\ p_3 &= 0.1, & \mathbf{z}_3 &= [1, 0, 0, 0, 1, 0, 0, 1, 0] \end{aligned}$$

$$\Rightarrow \boldsymbol{\mu} = [.3, 0, .7, 0, .3, .7, .7, .1, .2].$$

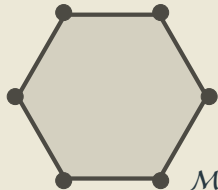


# Unstructured vs Structured

- Unstructured case: simplex  $\Delta$

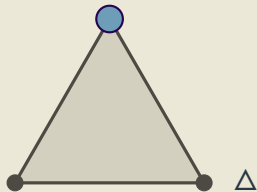


- Structured case: marginal polytope  $\mathcal{M}$

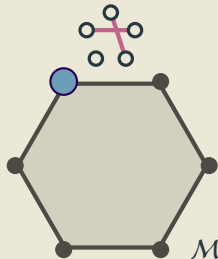


# Unstructured vs Structured

- Unstructured case: simplex  $\Delta$

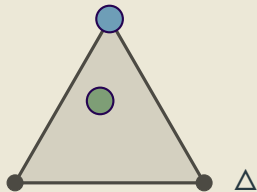


- Structured case: marginal polytope  $\mathcal{M}$

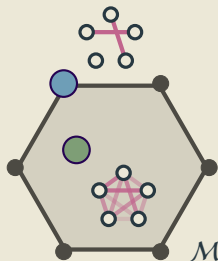


# Unstructured vs Structured

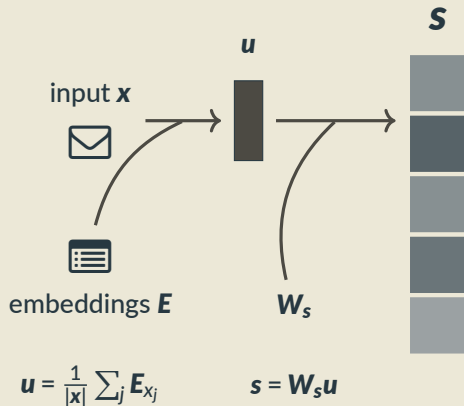
- Unstructured case: simplex  $\Delta$



- Structured case: marginal polytope  $\mathcal{M}$

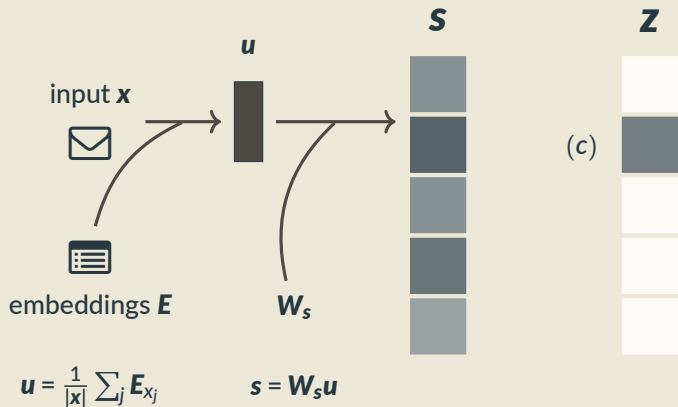


# Example: Regression with latent categorization



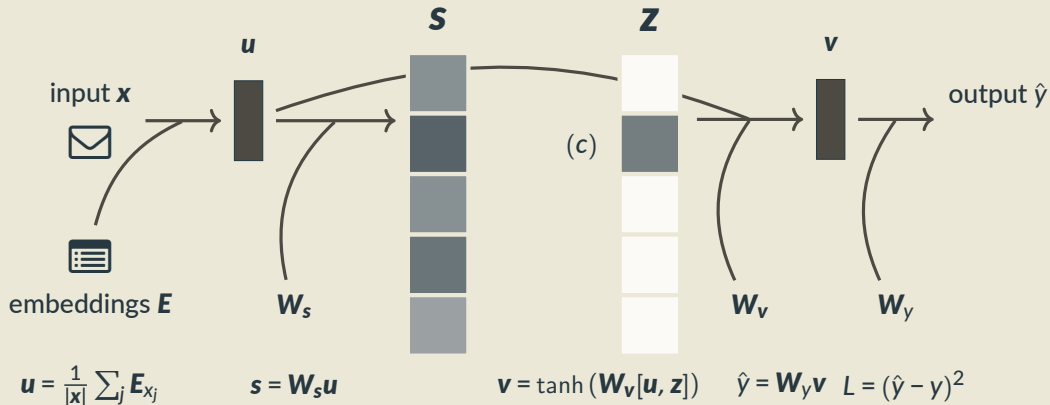


# Example: Regression with latent categorization




predict topic  $c$  ( $z = e_c$ )

# Example: Regression with latent categorization




predict topic  $c$  ( $z = e_c$ )

# Example: Regression with latent categorization




# Example: Regression with latent categorization

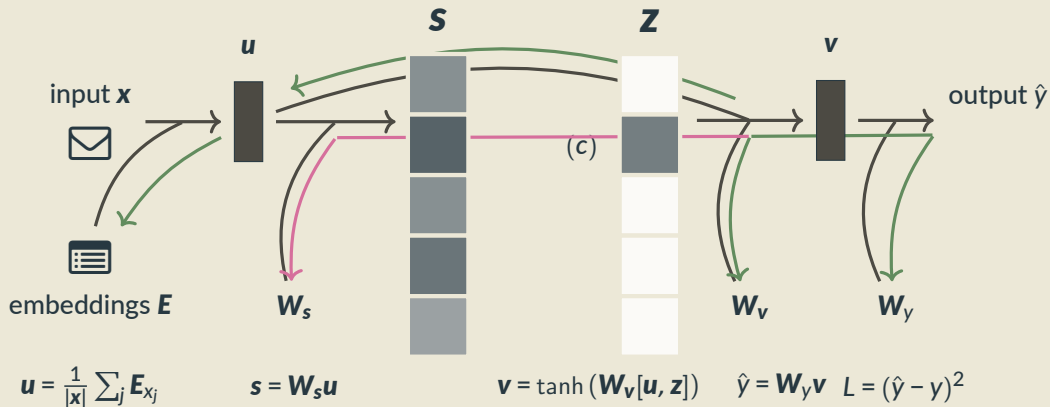


$$\frac{\partial L}{\partial W_s} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial v} \frac{\partial v}{\partial z} \underbrace{\frac{\partial z}{\partial s}}_{\equiv 0} \frac{\partial s}{\partial W_s}$$

# Example: Regression with latent categorization




# Example: Regression with latent categorization




Option 1. Pretrain latent classifier  $W_s$

# Example: Regression with latent categorization



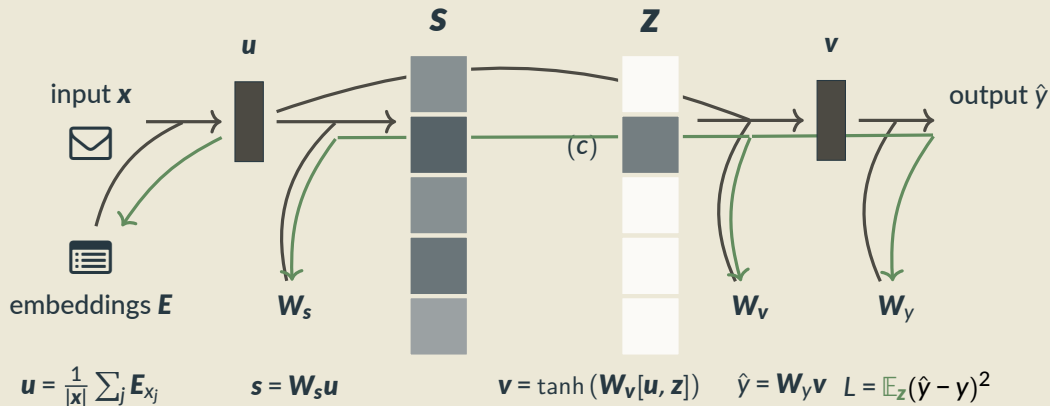
Option 2. Multi-task learning

# Example: Regression with latent categorization



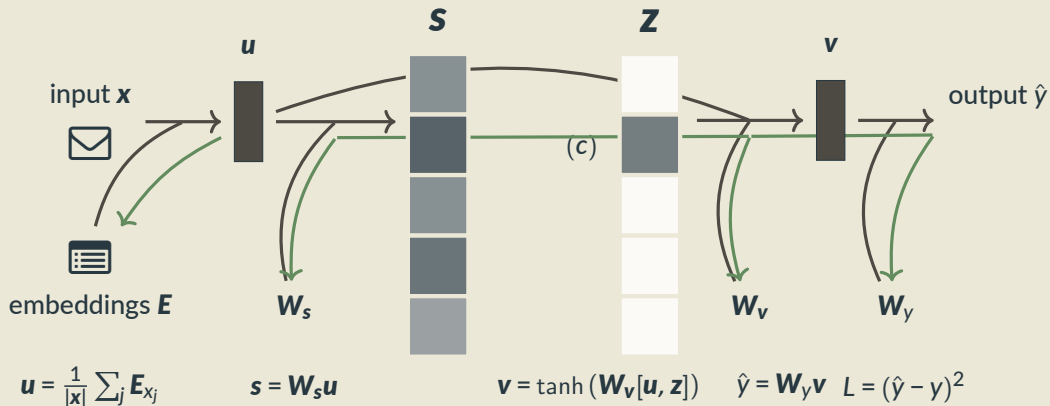


# Example: Regression with latent categorization



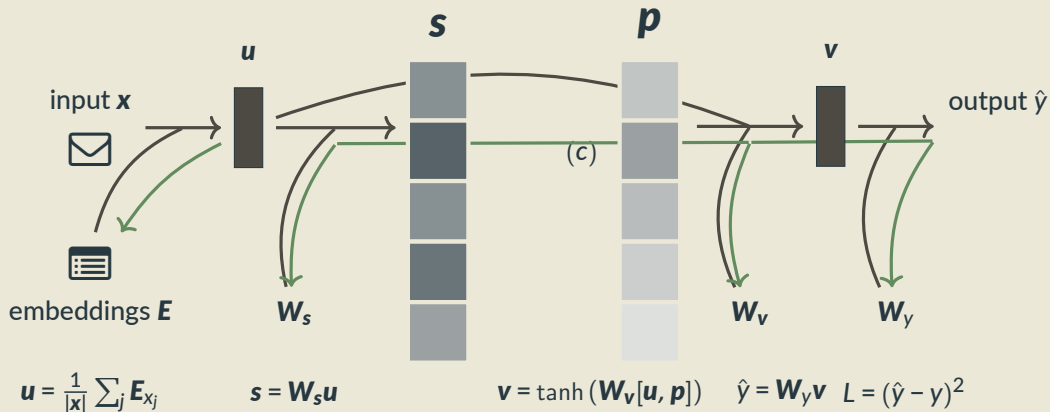
Option 3. Stochasticity!  $\frac{\partial \mathbb{E}_z (\hat{y}(z) - y)^2}{\partial W_s} \neq 0$

# Example: Regression with latent categorization



Option 4. Gradient surrogates (e.g. straight-through,  $\frac{\partial z}{\partial s} \leftarrow I$ )

# Example: Regression with latent categorization



Option 5. Continuous relaxation (e.g. softmax)

# Dealing with discrete latent variables

1. Pre-train external classifier
2. Multi-task learning
3. Stochastic latent variables
4. Gradient surrogates
5. Continuous relaxation

# Dealing with discrete latent variables

1. Pre-train external classifier
2. Multi-task learning
3. Stochastic latent variables (Part 2)
4. Gradient surrogates (Part 3)
5. Continuous relaxation (Part 4)

# Roadmap of the tutorial

- Part 1: Introduction ✓
- Part 2: Reinforcement learning

*Coffee Break*

- Part 3: Gradient surrogates
- Part 4: End-to-end differentiable models (1/2)

*Coffee Break*

- Part 4: End-to-end differentiable models (2/2)
- Part 5: Conclusions

## **II. Reinforcement Learning Methods**

# Latent structure via marginalization

- Given a sentence-label pair  $(x, y)$  and its **known** parse tree  $z$ ,



# Latent structure via marginalization

- Given a sentence-label pair  $(x, y)$  and its **known** parse tree  $\mathbf{z}$ , we can make a prediction  $\hat{y}(\mathbf{z}; x)$

# Latent structure via marginalization

- Given a sentence-label pair  $(x, y)$  and its **known** parse tree  $\mathbf{z}$ , we can make a prediction  $\hat{y}(\mathbf{z}; x)$  and incur a loss,

$$L(\hat{y}(\mathbf{z}; x), y)$$

# Latent structure via marginalization

- Given a sentence-label pair  $(x, y)$  and its **known** parse tree  $\mathbf{z}$ , we can make a prediction  $\hat{y}(\mathbf{z}; x)$  and incur a loss,

$$L(\hat{y}(\mathbf{z}; x), y) \text{ or simply } L(\mathbf{z})$$

# Latent structure via marginalization

- Given a sentence-label pair  $(x, y)$  and its **known** parse tree  $\mathbf{z}$ , we can make a prediction  $\hat{y}(\mathbf{z}; x)$  and incur a loss,

$$L(\hat{y}(\mathbf{z}; x), y) \text{ or simply } L(\mathbf{z})$$

- But we don't know  $\mathbf{z}$ !

# Latent structure via marginalization

- Given a sentence-label pair  $(x, y)$  and its **known** parse tree  $\mathbf{z}$ , we can make a prediction  $\hat{y}(\mathbf{z}; x)$  and incur a loss,

$$L(\hat{y}(\mathbf{z}; x), y) \text{ or simply } L(\mathbf{z})$$

- But we don't know  $\mathbf{z}$ !
- In this section:  
we jointly learn a structured prediction model  $\pi_{\theta}(\mathbf{z} \mid x)$

# Latent structure via marginalization

- Given a sentence-label pair  $(x, y)$  and its **known** parse tree  $\mathbf{z}$ , we can make a prediction  $\hat{y}(\mathbf{z}; x)$  and incur a loss,

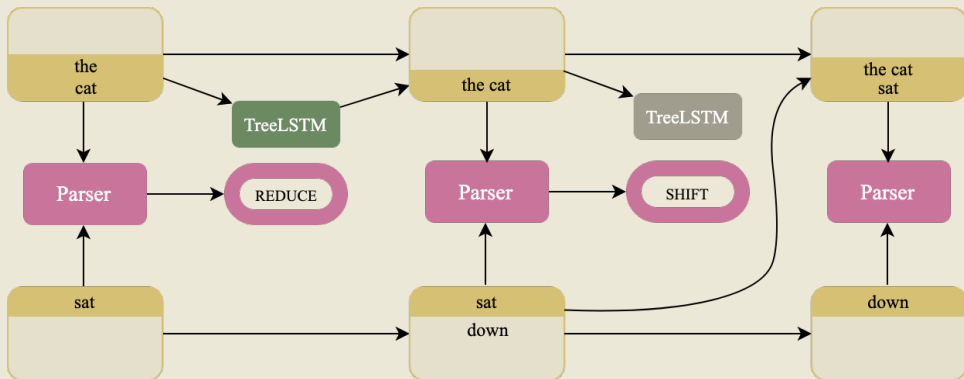
$$L(\hat{y}(\mathbf{z}; x), y) \text{ or simply } L(\mathbf{z})$$

- But we don't know  $\mathbf{z}$ !
- In this section:  
we jointly learn a structured prediction model  $\pi_{\theta}(\mathbf{z} \mid x)$  by optimizing the **expected loss**,

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z} \mid x)}[L(\mathbf{z})]$$

**But first, supervised  
SPINN**

# Stack-augmented Parser-Interpreter Neural-Network





# Stack-augmented Parser-Interpreter Neural-Network

- Joint learning: Combines a constituency parser and a sentence representation model.


# Stack-augmented Parser-Interpreter Neural-Network

- Joint learning: Combines a constituency parser and a sentence representation model.
- The parser,  $f_{\theta}(x)$  is a transition-based **shift-reduce** parser. It looks at top two elements of stack and top element of the buffer.


# Stack-augmented Parser-Interpreter Neural-Network

- Joint learning: Combines a constituency parser and a sentence representation model.
- The parser,  $f_{\theta}(x)$  is a transition-based **shift-reduce** parser. It looks at top two elements of stack and top element of the buffer.
- **TreeLSTM** combines top two elements of the stack when the parser chooses the REDUCE action.

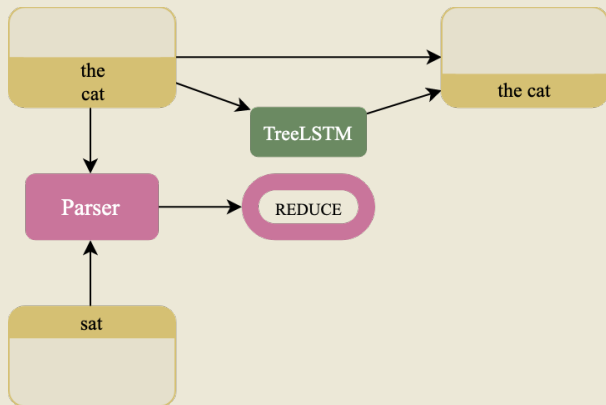
# Stack-augmented Parser-Interpreter Neural-Network



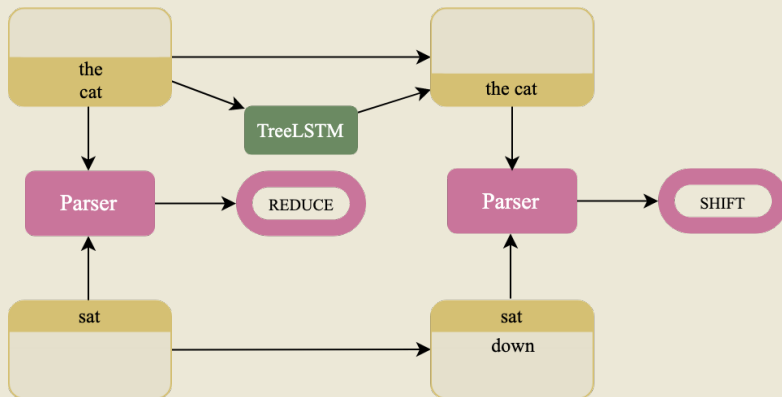
# Stack-augmented Parser-Interpreter Neural-Network



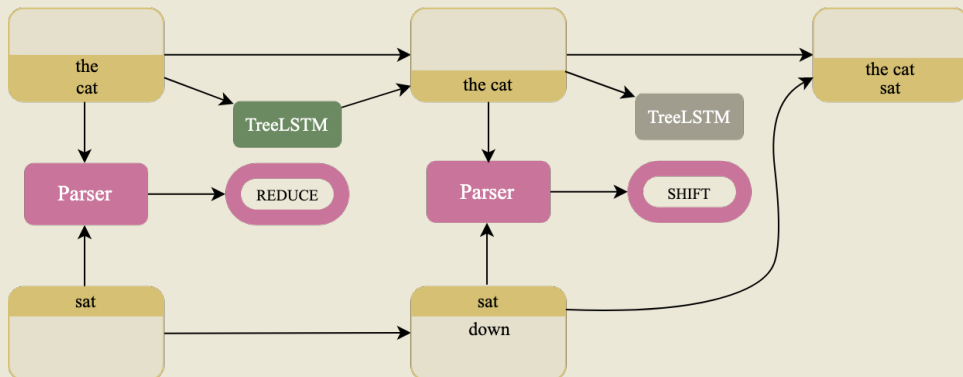
# Stack-augmented Parser-Interpreter Neural-Network



# Stack-augmented Parser-Interpreter Neural-Network

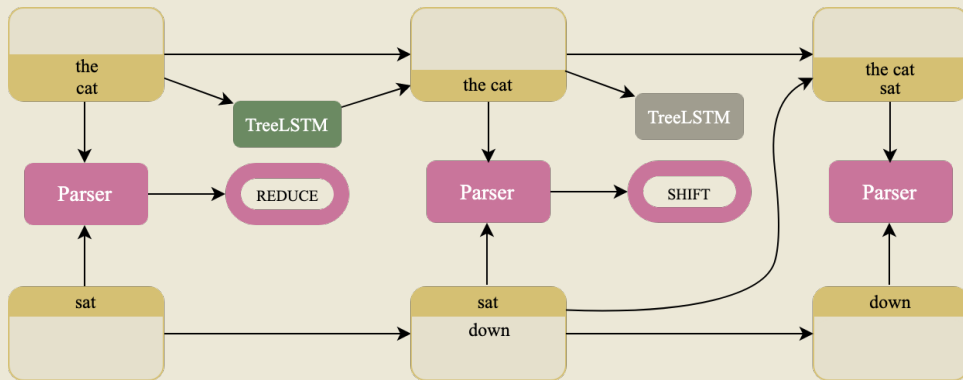


# Stack-augmented Parser-Interpreter Neural-Network





# Stack-augmented Parser-Interpreter Neural-Network



# Shift-Reduce parsing

We can write a shift-reduce style parse as a sequence of Bernoulli random variables,

$$\mathbf{z} = \{z_1, \dots, z_{2L-1}\}$$

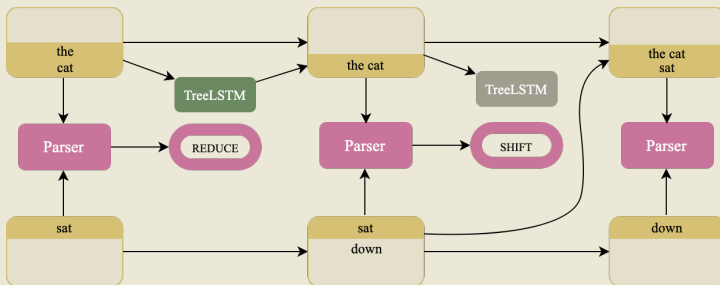
where,  $z_j \in \{0, 1\} \ \forall j \in [1, 2L - 1]$

# Shift-Reduce parsing

A sequence of Bernoulli trials but with conditional dependence,

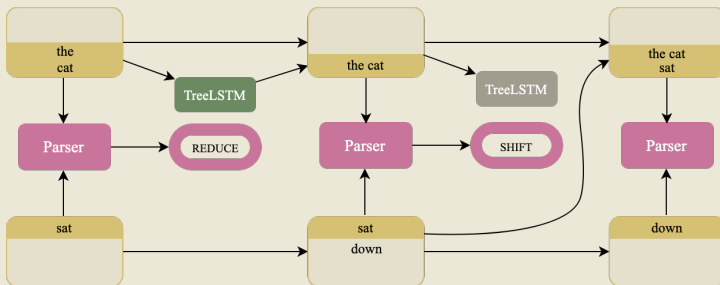
$$p(z_1, z_2, \dots, z_{2L-1}) = \prod_{j=1}^{2L-1} p(z_j \mid z_{<j})$$

# Latent structure learning with SPINN



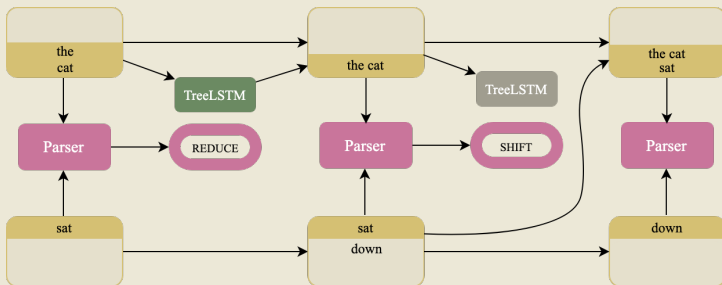
# Latent structure learning with SPINN

- But now, remove syntactic supervision from SPINN.



# Latent structure learning with SPINN

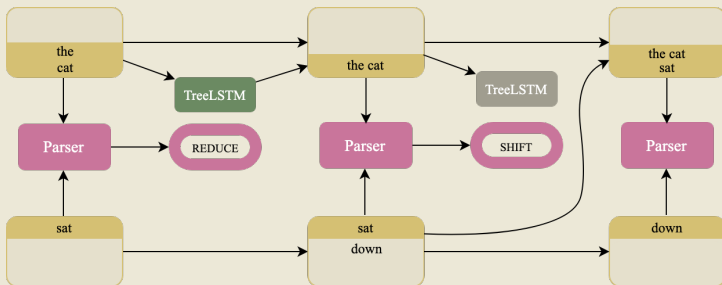
- But now, remove syntactic supervision from SPINN.



- We model the parse,  $\mathbf{z}$ , as a latent variable scored by  $f_{\theta}(\mathbf{x})$ .

# Latent structure learning with SPINN

- But now, remove syntactic supervision from SPINN.



- We model the parse,  $\mathbf{z}$ , as a latent variable scored by  $f_{\theta}(\mathbf{x})$ .
- With shift-reduce parsing, we're making discrete decisions  
⇒ REINFORCE as a “natural” solution.

# Unsupervised SPINN



# Unsupervised SPINN

No syntactic supervision.

Only reward is from the downstream task.

We only get this reward after parsing the full sentence.

# SPINN with REINFORCE

Some basic terminology,

- The action space is  $z_j \in \{\text{SHIFT}, \text{REDUCE}\}$ , and  $\mathbf{z}$  is a sequence of actions.

# SPINN with REINFORCE

Some basic terminology,

- The action space is  $z_j \in \{\text{SHIFT}, \text{REDUCE}\}$ , and  $\mathbf{z}$  is a sequence of actions.
- Training parser network parameters,  $\theta$  with REINFORCE

# SPINN with REINFORCE

## Some basic terminology,

- The action space is  $z_j \in \{\text{SHIFT}, \text{REDUCE}\}$ , and  $\mathbf{z}$  is a sequence of actions.
- Training parser network parameters,  $\theta$  with REINFORCE
- The state,  $\mathbf{h}$ , is the top two elements of the stack and the top element of the buffer.

# SPINN with REINFORCE

## Some basic terminology,

- The action space is  $z_j \in \{\text{SHIFT}, \text{REDUCE}\}$ , and  $\mathbf{z}$  is a sequence of actions.
- Training parser network parameters,  $\boldsymbol{\theta}$  with REINFORCE
- The state,  $\mathbf{h}$ , is the top two elements of the stack and the top element of the buffer.
- Learning a policy network  $\pi(\mathbf{z} \mid \mathbf{h}; \boldsymbol{\theta})$

# SPINN with REINFORCE

## Some basic terminology,

- The action space is  $z_j \in \{\text{SHIFT}, \text{REDUCE}\}$ , and  $\mathbf{z}$  is a sequence of actions.
- Training parser network parameters,  $\boldsymbol{\theta}$  with REINFORCE
- The state,  $\mathbf{h}$ , is the top two elements of the stack and the top element of the buffer.
- Learning a policy network  $\pi(\mathbf{z} \mid \mathbf{h}; \boldsymbol{\theta})$
- Maximize the reward, where  $\mathcal{R}$  is performance on the downstream task like sentence classification.

# SPINN with REINFORCE

## Some basic terminology,

- The action space is  $z_j \in \{\text{SHIFT}, \text{REDUCE}\}$ , and  $\mathbf{z}$  is a sequence of actions.
  - Training parser network parameters,  $\theta$  with REINFORCE
  - The state  $\mathbf{h}$  is the top two elements of the stack and the top element of the buffer
  - Learn
  - Max
- NOTE: Only a single reward at the end of parsing.
- ke  
sentence classification.

# Through the looking glass of REINFORCE

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{z} \sim \pi_{\boldsymbol{\theta}}(\mathbf{z}|\mathbf{x})} [L(\mathbf{z})]$$



# Through the looking glass of REINFORCE

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{z} \sim \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)} [L(\mathbf{z})] = \nabla_{\boldsymbol{\theta}} \left[ \sum_{\mathbf{z}} L(\mathbf{z}) \pi_{\boldsymbol{\theta}}(\mathbf{z} | x) \right]$$

*(Need to turn it into  $\mathbb{E}[\cdot]$  so we can MC-estimate)*

# Through the looking glass of REINFORCE

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{z} \sim \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)} [L(\mathbf{z})] = \nabla_{\boldsymbol{\theta}} \left[ \sum_{\mathbf{z}} L(\mathbf{z}) \pi_{\boldsymbol{\theta}}(\mathbf{z} | x) \right]$$

*(Need to turn it into  $\mathbb{E}[\cdot]$  so we can MC-estimate)*

$$= \sum_{\mathbf{z}} L(\mathbf{z}) \nabla_{\boldsymbol{\theta}} \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)$$

# Through the looking glass of REINFORCE

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{z} \sim \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)} [L(\mathbf{z})] = \nabla_{\boldsymbol{\theta}} \left[ \sum_{\mathbf{z}} L(\mathbf{z}) \pi_{\boldsymbol{\theta}}(\mathbf{z} | x) \right]$$

*(Need to turn it into  $\mathbb{E}[\cdot]$  so we can MC-estimate)*

$$= \sum_{\mathbf{z}} L(\mathbf{z}) \nabla_{\boldsymbol{\theta}} \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)$$

$$\nabla \log f = \frac{\nabla f}{f}, \text{ so } \nabla f = f \nabla \log f.$$

# Through the looking glass of REINFORCE

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{z} \sim \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)} [L(\mathbf{z})] = \nabla_{\boldsymbol{\theta}} \left[ \sum_{\mathbf{z}} L(\mathbf{z}) \pi_{\boldsymbol{\theta}}(\mathbf{z} | x) \right]$$

*(Need to turn it into  $\mathbb{E}[\cdot]$  so we can MC-estimate)*

$$= \sum_{\mathbf{z}} L(\mathbf{z}) \nabla_{\boldsymbol{\theta}} \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)$$

$$= \sum_{\mathbf{z}} L(\mathbf{z}) \pi_{\boldsymbol{\theta}}(\mathbf{z} | x) \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)$$

# Through the looking glass of REINFORCE

$$\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{z} \sim \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)} [L(\mathbf{z})] = \nabla_{\boldsymbol{\theta}} \left[ \sum_{\mathbf{z}} L(\mathbf{z}) \pi_{\boldsymbol{\theta}}(\mathbf{z} | x) \right]$$

*(Need to turn it into  $\mathbb{E}[\cdot]$  so we can MC-estimate)*

$$\begin{aligned} &= \sum_{\mathbf{z}} L(\mathbf{z}) \nabla_{\boldsymbol{\theta}} \pi_{\boldsymbol{\theta}}(\mathbf{z} | x) \\ &= \sum_{\mathbf{z}} L(\mathbf{z}) \pi_{\boldsymbol{\theta}}(\mathbf{z} | x) \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(\mathbf{z} | x) \\ &= \mathbb{E}_{\mathbf{z} \sim \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)} [L(\mathbf{z}) \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)] \end{aligned}$$

# SPINN with REINFORCE, aka RL-SPINN

Yogatama et al. [2017] uses REINFORCE to train SPINN!

# SPINN with REINFORCE, aka RL-SPINN

Yogatama et al. [2017] uses REINFORCE to train SPINN!

However, this vanilla implementation isn't very effective at learning syntax.

# SPINN with REINFORCE, aka RL-SPINN

Yogatama et al. [2017] uses REINFORCE to train SPINN!

However, this vanilla implementation isn't very effective at learning syntax.

This model fails to solve a simple toy problem.



# Toy problem: ListOps



# Toy problem: ListOps

Model	Accuracy
LSTM	74.4
RL-SPINN	64.8
TreeLSTM with ground-truth trees	98.7

# Toy problem: ListOps

Model	Accuracy
LSTM	74.4
RL-SPIN	4.8
TreeLST	8.7

But why?

# RL-SPINN's Troubles

This system faces at least two big problems,

# RL-SPINN's Troubles

This system faces at least two big problems,

1. High variance of gradients
2. Coadaptation

# High variance

- We have a single reward at the end of parsing.

# High variance

- We have a single reward at the end of parsing.
- We are sampling parses from very large search space!  
**Catalan number** of binary trees.

# High variance

- We have a single reward at the end of parsing.
- We are sampling parses from very large search space!

**Catalan number** of binary trees.

3 tokens  $\Rightarrow$  5 trees

5 tokens  $\Rightarrow$  42 trees

10 tokens  $\Rightarrow$  16796 trees



# High variance

- We have a single reward at the end of parsing.
- We are sampling parses from very large search space!  
**Catalan number** of binary trees.
- And the policy is stochastic.

# High variance

So, sometimes the policy lands in a “rewarding state”:




Figure: Truth: 7; Pred: 7

# High variance

Sometimes it doesn't:



Figure: Truth: 6; Pred: 5

# High variance

**Catalan number** of parses means we need many many samples to lower variance!

# High variance

**Catalan number** of parses means we need many many samples to lower variance!

Possible solutions:

1. Gradient normalization
2. Control variates, aka baselines

# Control variates

- A simple control variate: moving average of recent rewards

# Control variates

- A simple control variate: moving average of recent rewards
- Parameters are updated using the advantage which is the difference between the reward,  $\mathcal{R}$ , and the baseline prediction.

# Control variates

- A simple control variate: moving average of recent rewards
- Parameters are updated using the advantage which is the difference between the reward,  $\mathcal{R}$ , and the baseline prediction.

So,

$$\nabla \mathbb{E}_{\mathbf{z} \sim \pi(\mathbf{z})} = \mathbb{E}_{\mathbf{z} \sim \pi(\mathbf{z})} [(L(\mathbf{z}) - b(\mathbf{x})) \nabla \pi(\mathbf{z})]$$



# Issues with SPINN with REINFORCE

This system faces two big problems,

1. High variance of gradients
2. Coadaptation

# Coadaptation problem

Learning composition function parameters  $\phi$  with backpropagation,  
and parser parameters  $\theta$  with REINFORCE.

# Coadaptation problem

Learning composition function parameters  $\phi$  with backpropagation, and parser parameters  $\theta$  with REINFORCE.

Generally,  $\phi$  will be learned more quickly than  $\theta$ , making it harder to explore the parsing search space and optimize for  $\theta$ .

# Coadaptation problem

Learning composition function parameters  $\phi$  with backpropagation, and parser parameters  $\theta$  with REINFORCE.

Generally,  $\phi$  will be learned more quickly than  $\theta$ , making it harder to explore the parsing search space and optimize for  $\theta$ .

Difference in variance of two gradient estimates.

# Coadaptation problem

Learning composition function parameters  $\phi$  with backpropagation, and parser parameters  $\theta$  with REINFORCE.

Generally,  $\phi$  will be learned more quickly than  $\theta$ , making it harder to explore the parsing search space and optimize for  $\theta$ .

Difference in variance of two gradient estimates.

Possible solution: Proximal Policy Optimization [Schulman et al., 2017].

# Making REINFORCE+SPINN work

Havrylov et al. [2019] use,

1. Input dependent control variate
2. Gradient normalization
3. Proximal Policy Optimization

# Making REINFORCE+SPINN work

Havrylov et al. [2019] use,

1. Input dependent control variate
2. Gradient normalization
3. Proximal Policy Optimization

They solve ListOps!

# Making REINFORCE+SPINN work

Havrylov et al. [2019] use,

1. Input dependent control variate
2. Gradient normalization
3. Proximal Policy Optimization

They solve ListOps!

However, does not learn English grammars.



# Should I? Shouldn't I?

- Unbiased!

# Should I? Shouldn't I?

- Unbiased!
- High variance 🙄

# Should I? Shouldn't I?

- Unbiased!
- In a simple setting, with enough tricks, it can work! 😊
- High variance 😞

# Should I? Shouldn't I?

- Unbiased!
- In a simple setting, with enough tricks, it can work! 😊
- High variance 😞
- Has not yet been very effective at learning English syntax.

# Roadmap of the tutorial

- Part 1: Introduction ✓
- Part 2: Reinforcement learning ✓

*Coffee Break*

- Part 3: Gradient surrogates
- Part 4: End-to-end differentiable models (1/2)

*Coffee Break*

- Part 4: End-to-end differentiable models (2/2)
- Part 5: Conclusions

# III. Gradient Surrogates

# So far:

- Tackled **expected loss** in a **stochastic computation graph**

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

# So far:

- Tackled **expected loss** in a **stochastic computation graph**

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- Optimized with the **REINFORCE** estimator.



## So far:

- Tackled **expected loss** in a **stochastic computation graph**

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- Optimized with the **REINFORCE** estimator.
- Struggled with variance & sampling.

## So far:

- Tackled **expected loss** in a **stochastic computation graph**

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- Optimized with the **REINFORCE** estimator.
- Struggled with variance & sampling.

## In this section:

- Consider the **deterministic alternative**:

## So far:

- Tackled **expected loss** in a **stochastic computation graph**

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- Optimized with the **REINFORCE** estimator.
- Struggled with variance & sampling.

## In this section:

- Consider the **deterministic alternative**:

pick highest-score structure       $\hat{\mathbf{z}}(x) := \arg \max_{\mathbf{z} \in \mathcal{M}} \pi_{\theta}(\mathbf{z} | x)$

## So far:

- Tackled **expected loss** in a **stochastic computation graph**

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- Optimized with the **REINFORCE** estimator.
- Struggled with variance & sampling.

## In this section:

- Consider the **deterministic alternative**:

pick highest-score structure  
incur loss

$$\hat{\mathbf{z}}(x) := \arg \max_{\mathbf{z} \in \mathcal{M}} \pi_{\theta}(\mathbf{z} | x) \\ L(\hat{\mathbf{z}}(x))$$

## So far:

- Tackled **expected loss** in a **stochastic computation graph**

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- Optimized with the **REINFORCE** estimator.
- Struggled with variance & sampling.

## In this section:

- Consider the **deterministic alternative**:

pick highest-score structure

incur loss

$$\hat{\mathbf{z}}(x) := \arg \max_{\mathbf{z} \in \mathcal{M}} \pi_{\theta}(\mathbf{z} | x) \\ L(\hat{\mathbf{z}}(x))$$

- 3A: try to optimize the deterministic loss directly

## So far:

- Tackled **expected loss** in a **stochastic computation graph**

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- Optimized with the **REINFORCE** estimator.
- Struggled with variance & sampling.

## In this section:

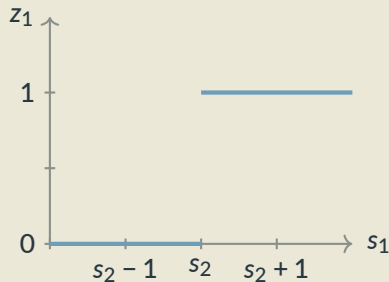
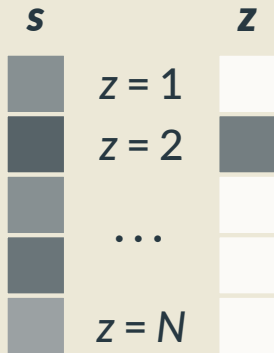
- Consider the **deterministic alternative**:

pick highest-score structure       $\hat{\mathbf{z}}(x) := \arg \max_{\mathbf{z} \in \mathcal{M}} \pi_{\theta}(\mathbf{z} | x)$   
incur loss       $L(\hat{\mathbf{z}}(x))$

- 3A: try to optimize the deterministic loss directly
- 3B: use this strategy to reduce variance in the stochastic model.

# Recap: The argmax problem

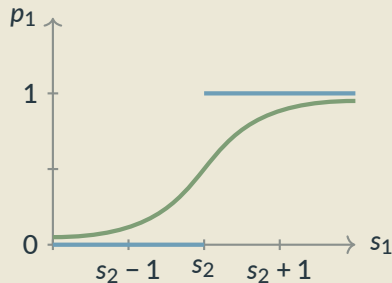
$$\mathbf{z} = \arg \max(\mathbf{s})$$



$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = \mathbf{0}$$

# Softmax

$$p_j = \exp(s_j)/Z$$



$$\frac{\partial \mathbf{p}}{\partial \mathbf{s}} = \text{diag}(\mathbf{p}) - \mathbf{p}\mathbf{p}^\top$$

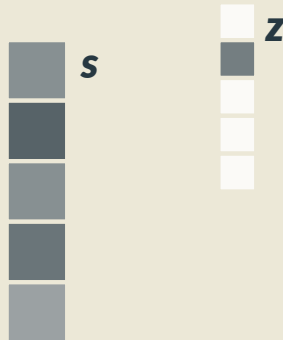


# Straight-Through Estimator



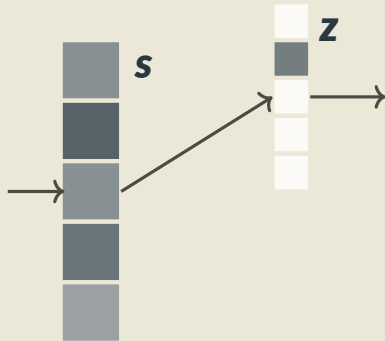
# Straight-Through Estimator

- Forward:  $\mathbf{z} = \arg \max(\mathbf{s})$



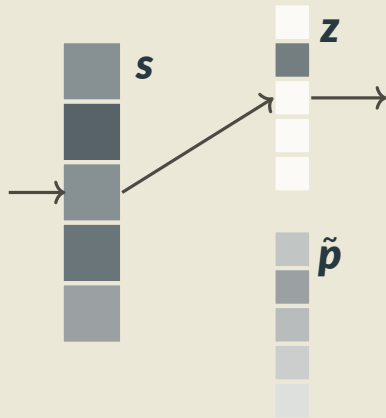
# Straight-Through Estimator

- Forward:  $\mathbf{z} = \arg \max(\mathbf{s})$



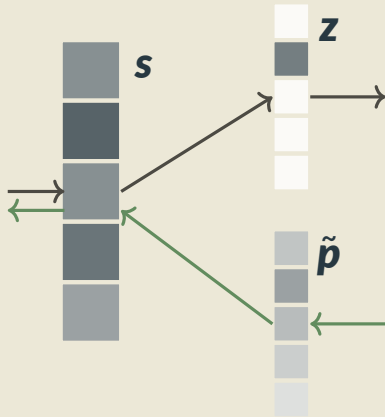
# Straight-Through Estimator

- Forward:  $\mathbf{z} = \arg \max(\mathbf{s})$
- Backward: pretend  $\mathbf{z}$  was some continuous  $\tilde{\mathbf{p}}$ ;  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} \neq \mathbf{0}$



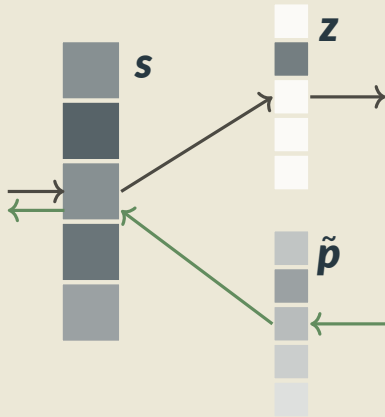
# Straight-Through Estimator

- Forward:  $\mathbf{z} = \arg \max(\mathbf{s})$
- Backward: pretend  $\mathbf{z}$  was some continuous  $\tilde{\mathbf{p}}$ ;  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} \neq \mathbf{0}$




# Straight-Through Estimator

- Forward:  $\mathbf{z} = \arg \max(\mathbf{s})$
- Backward: pretend  $\mathbf{z}$  was some continuous  $\tilde{\mathbf{p}}$ ;  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} \neq \mathbf{0}$ 
  - simplest: identity,  $\tilde{\mathbf{p}}(\mathbf{s}) = \mathbf{s}$ ,  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} = \mathbf{I}$



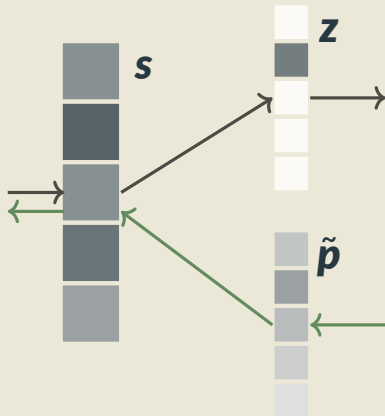
# Straight-Through Estimator

- Forward:  $\mathbf{z} = \arg \max(\mathbf{s})$
- Backward: pretend  $\mathbf{z}$  was some continuous  $\tilde{\mathbf{p}}$ ;  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} \neq \mathbf{0}$ 
  - simplest: identity,  $\tilde{\mathbf{p}}(\mathbf{s}) = \mathbf{s}$ ,  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} = \mathbf{I}$
  - others, e.g. softmax  $\tilde{\mathbf{p}}(\mathbf{s}) = \text{softmax}(\mathbf{s})$ ,  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} = \text{diag}(\tilde{\mathbf{p}}) - \tilde{\mathbf{p}}\tilde{\mathbf{p}}^T$



# Straight-Through Estimator

- Forward:  $\mathbf{z} = \arg \max(\mathbf{s})$
- Backward: pretend  $\mathbf{z}$  was some continuous  $\tilde{\mathbf{p}}$ ;  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} \neq \mathbf{0}$ 
  - simplest: identity,  $\tilde{\mathbf{p}}(\mathbf{s}) = \mathbf{s}$ ,  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} = \mathbf{I}$
  - others, e.g. softmax  $\tilde{\mathbf{p}}(\mathbf{s}) = \text{softmax}(\mathbf{s})$ ,  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} = \text{diag}(\tilde{\mathbf{p}}) - \tilde{\mathbf{p}}\tilde{\mathbf{p}}^T$
- More explanation in a while

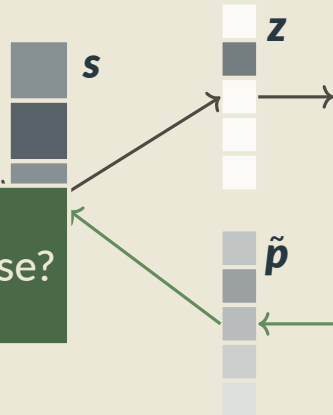




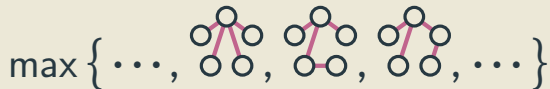
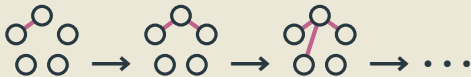
# Straight-Through Estimator

- Forward:  $\mathbf{z} = \arg \max(\mathbf{s})$
- Backward: pretend  $\mathbf{z}$  was some continuous  $\tilde{\mathbf{p}}$ ;  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} \neq \mathbf{0}$ 
  - simplest: identity,  $\tilde{\mathbf{p}}(\mathbf{s}) = \mathbf{s}$ ,  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} = \mathbf{I}$
  - others, e.g. softmax  $\tilde{\mathbf{p}}(\mathbf{s}) = \text{softmax}(\mathbf{s})$ ,  $\frac{\partial \tilde{\mathbf{p}}}{\partial \mathbf{s}} = \text{diag}(\tilde{\mathbf{p}}) - \tilde{\mathbf{p}}\tilde{\mathbf{p}}^T$
- More explanation

What about the structured case?



# Dealing with the combinatorial explosion



## 1. Incremental structures

- Build structure **greedily**, as sequence of discrete choices (e.g., shift-reduce).
- Scores (partial structure, action) tuples.
- **Advantages:** flexible, rich histories.
- **Disadvantages:** greedy, local decisions are suboptimal, error propagation.

## 2. Factorization into parts

- Optimizes **globally** (e.g. Viterbi, Chu-Liu-Edmonds, Kuhn-Munkres).
- Scores smaller parts.
- **Advantages:** optimal, elegant, can handle hard & global constraints.
- **Disadvantages:** strong assumptions.

# STE for incremental structures

# STE for incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)

# STE for incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)
- Assigns a score to any (partial structure, action) tuple.

# STE for incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)
- Assigns a score to any (partial structure, action) tuple.
- In this case, we just apply the straight-through estimator for each step.

# STE for incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)
- Assigns a score to any (partial structure, action) tuple.
- In this case, we just apply the straight-through estimator for each step.
- Forward: the **highest scoring action** for each step

# STE for incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)
- Assigns a score to any (partial structure, action) tuple.
- In this case, we just apply the straight-through estimator for each step.
- Forward: the **highest scoring action** for each step
- Backward: pretend that we had used a **differentiable surrogate function**



# STE for incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)
- Assigns a score to any (partial structure, action) tuple.
- In this case, we just apply the straight-through estimator for each step.
- Forward: the **highest scoring action** for each step
- Backward: pretend that we had used a **differentiable surrogate function**

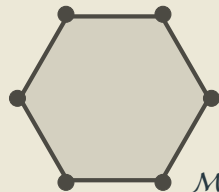
Example: Latent Tree Learning with Differentiable Parsers: Shift-Reduce Parsing and Chart Parsing [Maillard and Clark, 2018] (STE through beam search).

# STE for the factorized approach

Requires a bit more work:


- Recap: marginal polytope
- Predicting structures globally: Maximum A Posteriori (MAP)
- Deriving Straight-Through and SPIGOT

# The structured case: Marginal polytope



# The structured case: Marginal polytope

- Each vertex corresponds to one such *bit* vector  $\mathbf{z}$



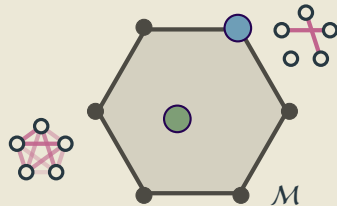
# The structured case: Marginal polytope

- Each vertex corresponds to one such *bit vector*  $\mathbf{z}$
- Points inside correspond to *marginal distributions*: convex combinations of structured objects

$$\boldsymbol{\mu} = \underbrace{p_1 \mathbf{z}_1 + \dots + p_N \mathbf{z}_N}_{\text{exponentially many terms}}, \quad \mathbf{p} \in \Delta.$$

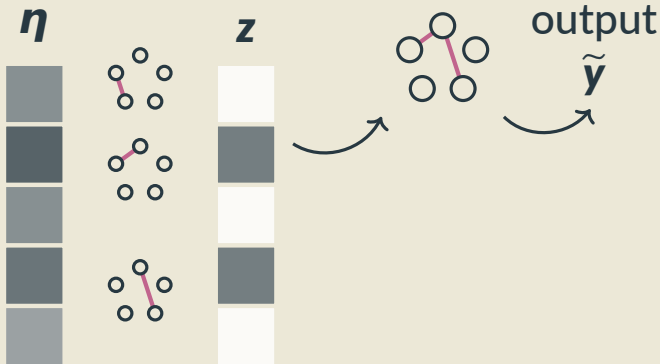
$$\begin{aligned} p_1 &= 0.2, & \mathbf{z}_1 &= [1, 0, 0, 0, 1, 0, 0, 0, 1] \\ p_2 &= 0.7, & \mathbf{z}_2 &= [0, 0, 1, 0, 0, 1, 1, 0, 0] \\ p_3 &= 0.1, & \mathbf{z}_3 &= [1, 0, 0, 0, 1, 0, 0, 1, 0] \end{aligned}$$

$$\Rightarrow \boldsymbol{\mu} = [.3, 0, .7, 0, .3, .7, .7, .1, .2].$$



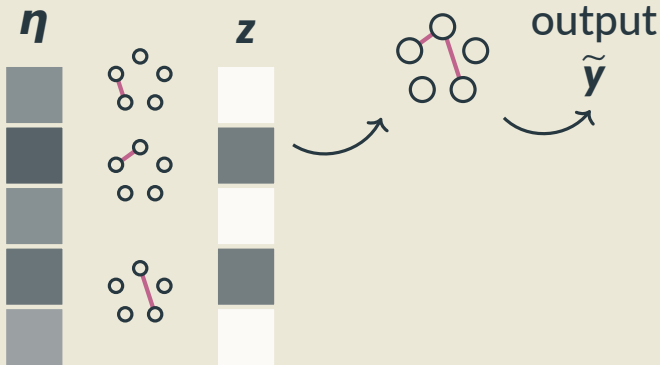
# Predicting structures from scores of parts

- $\eta(i \rightarrow j)$ : score of arc  $i \rightarrow j$
- $z(i \rightarrow j)$ : is arc  $i \rightarrow j$  selected?



# Predicting structures from scores of parts

- $\eta(i \rightarrow j)$ : score of arc  $i \rightarrow j$
- $z(i \rightarrow j)$ : is arc  $i \rightarrow j$  selected?
- Task-specific algorithm for the highest-scoring structure.



# Algorithms for specific structures

## Best structure (MAP)

**Sequence tagging**

Viterbi  
[Rabiner, 1989]

**Constituent trees**

CKY  
[Kasami, 1966, Younger, 1967]  
[Cocke and Schwartz, 1970]

**Temporal alignments**

DTW  
[Sakoe and Chiba, 1978]

**Dependency trees**

Max. Spanning Arborescence  
[Chu and Liu, 1965, Edmonds, 1967]


**Assignments**

Kuhn-Munkres  
[Kuhn, 1955, Jonker and Volgenant, 1987]



# Structured Straight-Through

- Forward pass:  
Find highest-scoring structure:  
 $\mathbf{z} = \arg \max_{\mathbf{z} \in \mathcal{Z}} \boldsymbol{\eta}^\top \mathbf{z}$
- Backward pass:  
pretend we used  $\tilde{\boldsymbol{\mu}} = \boldsymbol{\eta}$ .



# Straight-Through Estimator

## Revisited

# Straight-Through Estimator

## Revisited

- In the forward pass,  $\mathbf{z} = \arg \max(\mathbf{s})$ .

# Straight-Through Estimator

## Revisited

- In the forward pass,  $\mathbf{z} = \arg \max(\mathbf{s})$ .
- if we had labels (multi-task learning),  $L_{\text{MTL}} = L(\hat{y}(\mathbf{z}), y) + L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}})$

# Straight-Through Estimator

## Revisited

- In the forward pass,  $\mathbf{z} = \arg \max(\mathbf{s})$ .
- if we had labels (multi-task learning),  $L_{\text{MTL}} = L(\hat{y}(\mathbf{z}), y) + L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}})$
- One choice: perceptron loss  $L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}}) = \mathbf{s}^T \mathbf{z} - \mathbf{s}^T \mathbf{z}^{\text{true}}$ ;  $\frac{\partial L_{\text{hid}}}{\partial \mathbf{s}} = \mathbf{z} - \mathbf{z}^{\text{true}}$ .

# Straight-Through Estimator

## Revisited

- In the forward pass,  $\mathbf{z} = \arg \max(\mathbf{s})$ .
- if we had labels (multi-task learning),  $L_{\text{MTL}} = L(\hat{y}(\mathbf{z}), y) + L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}})$
- One choice: perceptron loss  $L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}}) = \mathbf{s}^T \mathbf{z} - \mathbf{s}^T \mathbf{z}^{\text{true}}$ ;  $\frac{\partial L_{\text{hid}}}{\partial \mathbf{s}} = \mathbf{z} - \mathbf{z}^{\text{true}}$ .
- We don't have labels! Induce labels by “pulling back” the downstream target:  
the “best” (unconstrained) latent value would be:  $\arg \min_{\tilde{\mathbf{z}} \in \mathbb{R}^D} L(\hat{y}(\tilde{\mathbf{z}}), y)$

# Straight-Through Estimator

## Revisited

- In the forward pass,  $\mathbf{z} = \arg \max(\mathbf{s})$ .
- if we had labels (multi-task learning),  $L_{\text{MTL}} = L(\hat{\mathbf{y}}(\mathbf{z}), y) + L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}})$
- One choice: perceptron loss  $L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}}) = \mathbf{s}^T \mathbf{z} - \mathbf{s}^T \mathbf{z}^{\text{true}}$ ;  $\frac{\partial L_{\text{hid}}}{\partial \mathbf{s}} = \mathbf{z} - \mathbf{z}^{\text{true}}$ .
- We don't have labels! Induce labels by “pulling back” the downstream target:  
the “best” (unconstrained) latent value would be:  $\arg \min_{\tilde{\mathbf{z}} \in \mathbb{R}^D} L(\hat{\mathbf{y}}(\tilde{\mathbf{z}}), y)$
- One gradient descent step starting from  $\mathbf{z}$ :  $\mathbf{z}^{\text{true}} \leftarrow \mathbf{z} - \frac{\partial L}{\partial \mathbf{z}}$

# Straight-Through Estimator

## Revisited

- In the forward pass,  $\mathbf{z} = \arg \max(\mathbf{s})$ .
- if we had labels (multi-task learning),  $L_{\text{MTL}} = L(\hat{y}(\mathbf{z}), y) + L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}})$
- One choice: perceptron loss  $L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}}) = \mathbf{s}^T \mathbf{z} - \mathbf{s}^T \mathbf{z}^{\text{true}}$ ;  $\frac{\partial L_{\text{hid}}}{\partial \mathbf{s}} = \mathbf{z} - \mathbf{z}^{\text{true}}$ .
- We don't have labels! Induce labels by “pulling back” the downstream target:  
the “best” (unconstrained) latent value would be:  $\arg \min_{\tilde{\mathbf{z}} \in \mathbb{R}^D} L(\hat{y}(\tilde{\mathbf{z}}), y)$
- One gradient descent step starting from  $\mathbf{z}$ :  $\mathbf{z}^{\text{true}} \leftarrow \mathbf{z} - \frac{\partial L}{\partial \mathbf{z}}$

$$\frac{\partial L_{\text{MTL}}}{\partial \mathbf{s}} = \underbrace{\frac{\partial L}{\partial \mathbf{s}}}_{=0} + \frac{\partial L_{\text{hid}}}{\partial \mathbf{s}}$$



# Straight-Through Estimator

## Revisited

- In the forward pass,  $\mathbf{z} = \arg \max(\mathbf{s})$ .
- if we had labels (multi-task learning),  $L_{\text{MTL}} = L(\hat{y}(\mathbf{z}), y) + L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}})$
- One choice: perceptron loss  $L_{\text{hid}}(\mathbf{s}, \mathbf{z}^{\text{true}}) = \mathbf{s}^T \mathbf{z} - \mathbf{s}^T \mathbf{z}^{\text{true}}$ ;  $\frac{\partial L_{\text{hid}}}{\partial \mathbf{s}} = \mathbf{z} - \mathbf{z}^{\text{true}}$ .
- We don't have labels! Induce labels by “pulling back” the downstream target:  
the “best” (unconstrained) latent value would be:  $\arg \min_{\tilde{\mathbf{z}} \in \mathbb{R}^D} L(\hat{y}(\tilde{\mathbf{z}}), y)$
- One gradient descent step starting from  $\mathbf{z}$ :  $\mathbf{z}^{\text{true}} \leftarrow \mathbf{z} - \frac{\partial L}{\partial \mathbf{z}}$

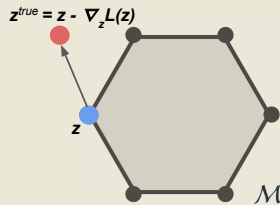
$$\frac{\partial L_{\text{MTL}}}{\partial \mathbf{s}} = \underbrace{\frac{\partial L}{\partial \mathbf{s}}}_{=0} + \frac{\partial L_{\text{hid}}}{\partial \mathbf{s}} = \mathbf{z} - \left( \mathbf{z} - \frac{\partial L}{\partial \mathbf{z}} \right) = \frac{\partial L}{\partial \mathbf{z}}$$

# Straight-Through in the structured case

- Structured STE: perceptron update with induced annotation

$$\arg \min_{\boldsymbol{\mu} \in \mathbb{R}^D} L(\hat{y}(\boldsymbol{\mu}), y) \approx \mathbf{z} - \nabla_{\mathbf{z}} L(\mathbf{z}) \rightarrow \mathbf{z}^{\text{true}}$$

(one step of gradient descent)



# Straight-Through in the structured case

- Structured STE: perceptron update with induced annotation



$$\arg \min_{\boldsymbol{\mu} \in \mathbb{R}^D} L(\hat{y}(\boldsymbol{\mu}), y) \approx \mathbf{z} - \nabla_{\mathbf{z}} L(\mathbf{z}) \rightarrow \mathbf{z}^{\text{true}}$$

(one step of gradient descent)

- SPIGOT takes into account the constraints; uses the induced annotation

$$\arg \min_{\boldsymbol{\mu} \in \mathcal{M}} L(\hat{y}(\boldsymbol{\mu}), y) \approx \text{Proj}_{\mathcal{M}}(\mathbf{z} - \nabla_{\mathbf{z}} L(\mathbf{z})) \rightarrow \mathbf{z}^{\text{true}}$$

(one step of *projected* gradient descent!)



# Straight-Through in the structured case

- Structured STE: perceptron update with induced annotation

$$\arg \min_{\boldsymbol{\mu} \in \mathbb{R}^D} L(\hat{y}(\boldsymbol{\mu}), y) \approx \mathbf{z} - \nabla_{\mathbf{z}} L(\mathbf{z}) \rightarrow \mathbf{z}^{\text{true}}$$

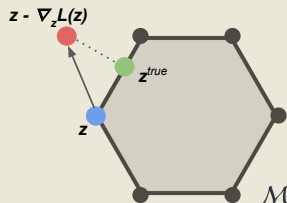

(one step of gradient descent)

- SPIGOT takes into account the constraints; uses the induced annotation

$$\arg \min_{\boldsymbol{\mu} \in \mathcal{M}} L(\hat{y}(\boldsymbol{\mu}), y) \approx \text{Proj}_{\mathcal{M}}(\mathbf{z} - \nabla_{\mathbf{z}} L(\mathbf{z})) \rightarrow \mathbf{z}^{\text{true}}$$

(one step of *projected* gradient descent!)

- We discuss a generic way to compute the projection in part 4.



# Summary: Straight-Through Estimator

# Summary: Straight-Through Estimator

We saw how to use the *Straight-Through Estimator* to allow learning models with *argmax* in the middle of the computation graph.

# Summary: Straight-Through Estimator

We saw how to use the *Straight-Through Estimator* to allow learning models with *argmax* in the middle of the computation graph.

We were optimizing  $L(\hat{\mathbf{z}}(x))$

# Summary: Straight-Through Estimator

We saw how to use the *Straight-Through Estimator* to allow learning models with *argmax* in the middle of the computation graph.

We were optimizing  $L(\hat{\mathbf{z}}(x))$

Now we will see how to apply STE for stochastic graphs, as an alternative approach of REINFORCE.



# Stochastic node in the computation graph

# Stochastic node in the computation graph

Recall the stochastic objective:

$$\mathbb{E}_{\pi_{\boldsymbol{\theta}}(\mathbf{z}|\mathbf{x})}[L(\mathbf{z})]$$

# Stochastic node in the computation graph

Recall the stochastic objective:

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- REINFORCE (previous section).

# Stochastic node in the computation graph

Recall the stochastic objective:

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- REINFORCE (previous section). High variance. 🙄

# Stochastic node in the computation graph

Recall the stochastic objective:

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- REINFORCE (previous section). High variance. 😞
- An alternative is using the *reparameterization trick* [Kingma and Welling, 2014].

# Categorical reparameterization

# Categorical reparameterization

- Sampling from a categorical value in the middle of the computation graph.

$$\mathbf{z} \sim \pi_{\theta}(\mathbf{z} \mid \mathbf{x}) \propto \exp \mathbf{s}_{\theta}(\mathbf{z} \mid \mathbf{x})$$

s



z



# Categorical reparameterization

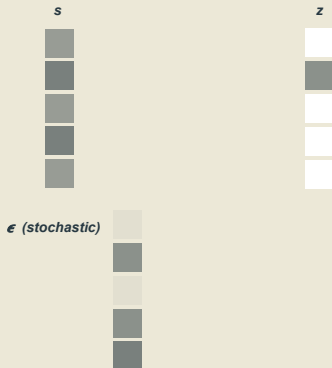
- Sampling from a categorical value in the middle of the computation graph.  
 $\mathbf{z} \sim \pi_{\boldsymbol{\theta}}(\mathbf{z} \mid \mathbf{x}) \propto \exp \mathbf{s}_{\boldsymbol{\theta}}(\mathbf{z} \mid \mathbf{x})$
- What is the gradient of a sample  $\frac{\partial \mathbf{z}}{\partial \boldsymbol{\theta}}$  ?!






# Categorical reparameterization

- Sampling from a categorical value in the middle of the computation graph.  
 $\mathbf{z} \sim \pi_{\theta}(\mathbf{z} | \mathbf{x}) \propto \exp \mathbf{s}_{\theta}(\mathbf{z} | \mathbf{x})$
- What is the gradient of a sample  $\frac{\partial \mathbf{z}}{\partial \theta}$  ?!
- Reparameterization: Move the stochasticity out of the gradient path.




# Categorical reparameterization

- Sampling from a categorical value in the middle of the computation graph.  
 $\mathbf{z} \sim \pi_{\boldsymbol{\theta}}(\mathbf{z} | \mathbf{x}) \propto \exp \mathbf{s}_{\boldsymbol{\theta}}(\mathbf{z} | \mathbf{x})$
- What is the gradient of a sample  $\frac{\partial \mathbf{z}}{\partial \boldsymbol{\theta}}$  ?!
- Reparameterization: Move the stochasticity out of the gradient path.
- Makes  $\mathbf{z}$  deterministic w.r.t.  $\mathbf{s}$ !



# Categorical reparameterization

- Sampling from a categorical value in the middle of the computation graph.  
 $\mathbf{z} \sim \pi_{\theta}(\mathbf{z} \mid \mathbf{x}) \propto \exp \mathbf{s}_{\theta}(\mathbf{z} \mid \mathbf{x})$
- What is the gradient of a sample  $\frac{\partial \mathbf{z}}{\partial \theta}$  ?!
- Reparameterization: Move the stochasticity out of the gradient path.
- Makes  $\mathbf{z}$  deterministic w.r.t.  $\mathbf{s}$ !



# Categorical reparameterization

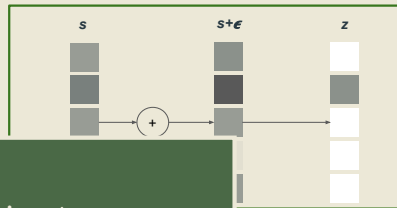
- Sampling from a categorical value in the middle of the computation graph.

$$\mathbf{z} \sim \pi_{\theta}(\mathbf{z} | \mathbf{x}) \propto \exp \mathbf{s}_{\theta}(\mathbf{z} | \mathbf{x})$$

- What is the gradient of the loss with respect to  $\mathbf{z}$ ?

- Reparameterization of stochasticity

- Makes  $\mathbf{z}$  deterministic



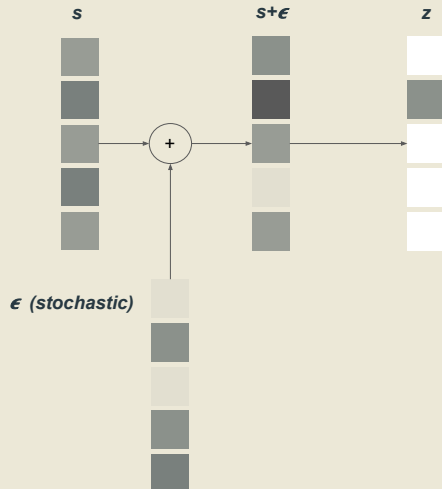
As a result:

Stochasticity is moved as an input.


We can backpropagate through the deterministic input to  $\mathbf{z}$ .



# Categorical reparameterization



# Categorical reparameterization



How do we sample from a categorical variable?



# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:



# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$
- return  $\mathbf{z} = \mathbf{e}_t$  s.t.  $c_t \leq u < c_{t+1}$

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:      2. The Gumbel-Max trick:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$
- return  $\mathbf{z} = \mathbf{e}_t$  s.t.  $c_t \leq u < c_{t+1}$

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$
- return  $\mathbf{z} = \mathbf{e}_t$  s.t.  $c_t \leq u < c_{t+1}$

## 2. The Gumbel-Max trick:

- $u_i \sim \text{Uniform}(0, 1)$

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$
- return  $\mathbf{z} = \mathbf{e}_t$  s.t.  $c_t \leq u < c_{t+1}$

## 2. The Gumbel-Max trick:

- $u_i \sim \text{Uniform}(0, 1)$
- $\epsilon_i = -\log(-\log(u_i))$

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$
- return  $\mathbf{z} = \mathbf{e}_t$  s.t.  $c_t \leq u < c_{t+1}$

## 2. The Gumbel-Max trick:

- $u_i \sim \text{Uniform}(0, 1)$
- $\epsilon_i = -\log(-\log(u_i))$
- $\mathbf{z} = \arg \max(\mathbf{s} + \boldsymbol{\epsilon})$



# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$
- return  $\mathbf{z} = \mathbf{e}_t$  s.t.  $c_t \leq u < c_{t+1}$

## 2. The Gumbel-Max trick:

- $u_i \sim \text{Uniform}(0, 1)$
- $\epsilon_i = -\log(-\log(u_i))$
- $\mathbf{z} = \arg \max(\mathbf{s} + \boldsymbol{\epsilon})$

The two methods are equivalent. (*Not obvious, but we will not prove it now.*)

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$
- return  $\mathbf{z} = \mathbf{e}_t$  s.t.  $c_t \leq u < c_{t+1}$

## 2. The Gumbel-Max trick:

- $u_i \sim \text{Uniform}(0, 1)$
- $\epsilon_i = -\log(-\log(u_i))$
- $\mathbf{z} = \arg \max(\mathbf{s} + \boldsymbol{\epsilon})$

The two methods are equivalent. (*Not obvious, but we will not prove it now.*)  
Requires sampling from the Standard Gumbel Distribution  $G(0,1)$ .

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$
- return  $\mathbf{z} = \mathbf{e}_t$  s.t.  $c_t \leq u < c_{t+1}$

## 2. The Gumbel-Max trick:

- $u_i \sim \text{Uniform}(0, 1)$
- $\epsilon_i = -\log(-\log(u_i))$
- $\mathbf{z} = \arg \max(\mathbf{s} + \boldsymbol{\epsilon})$

The two methods are equivalent. *(Not obvious, but we will not prove it now.)*

Requires sampling from the Standard Gumbel Distribution  $G(0,1)$ .

*Derivation & more info: [Adams, 2013, Vieira, 2014]*


# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling • 2 The Gumbel-Max trick:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$
- return  $\mathbf{z} = \mathbf{e}_t$

The t  
f



now.)  
).

# Sampling from a categorical variable

We want to sample from a categorical variable with scores  $\mathbf{s}$  (class  $i$  has a score  $s_i$ )

## 1. Inverse transform sampling • 2 The Gumbel-Max trick:

- $\mathbf{p} = \text{softmax}(\mathbf{s})$
- $c_i = \sum_{j \leq i} p_j$
- $u \sim \text{Uniform}(0, 1)$
- return  $\mathbf{z} = \mathbf{e}_t$

The t  
f



now.)  
).


# Straight-Through Gumbel Estimator

Apply a variant of the Straight-Through Estimator to Gumbel-Max!

# Straight-Through Gumbel Estimator

Apply a variant of the Straight-Through Estimator to Gumbel-Max!


- Forward:  $\mathbf{z} = \arg \max(\mathbf{s} + \boldsymbol{\epsilon})$



# Straight-Through Gumbel Estimator

Apply a variant of the Straight-Through Estimator to Gumbel-Max!

- Forward:  $\mathbf{z} = \arg \max(\mathbf{s} + \boldsymbol{\epsilon})$
- Backward: pretend we had done  
 $\tilde{\mathbf{p}} = \text{softmax}(\mathbf{s} + \boldsymbol{\epsilon})$






# Straight-Through Gumbel Estimator

Apply a variant of the Straight-Through Estimator to Gumbel-Max!

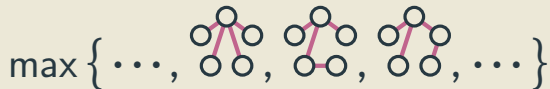
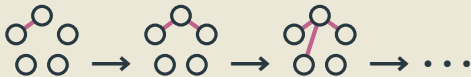
- Forward:  $\mathbf{z} = \arg \max(\mathbf{s} + \boldsymbol{\epsilon})$
- Backward: pretend we had done

$$\tilde{\mathbf{p}} = \text{softmax}(\mathbf{s} + \boldsymbol{\epsilon})$$

What about the structured case?



# Dealing with the combinatorial explosion



## 1. Incremental structures

- Build structure **greedily**, as sequence of discrete choices (e.g., shift-reduce).
- Scores (partial structure, action) tuples.
- **Advantages:** flexible, rich histories.
- **Disadvantages:** greedy, local decisions are suboptimal, error propagation.

## 2. Factorization into parts

- Optimizes **globally** (e.g. Viterbi, Chu-Liu-Edmonds, Kuhn-Munkres).
- Scores smaller parts.
- **Advantages:** optimal, elegant, can handle hard & global constraints.
- **Disadvantages:** strong assumptions.

# Sampling from incremental structures

# Sampling from incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)

# Sampling from incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)
- Assigns a score to any (partial structure, action) tuple.

# Sampling from incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)
- Assigns a score to any (partial structure, action) tuple.
- Reparameterize the scores with Gumbel-Max - now we have a deterministic node.

# Sampling from incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)
- Assigns a score to any (partial structure, action) tuple.
- Reparameterize the scores with Gumbel-Max - now we have a deterministic node.
- Forward: the **argmax** from the reparameterized scores for each step

# Sampling from incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)
- Assigns a score to any (partial structure, action) tuple.
- Reparameterize the scores with Gumbel-Max - now we have a deterministic node.
- Forward: the **argmax** from the reparameterized scores for each step
- Backward: pretend we had used a **differentiable surrogate function**



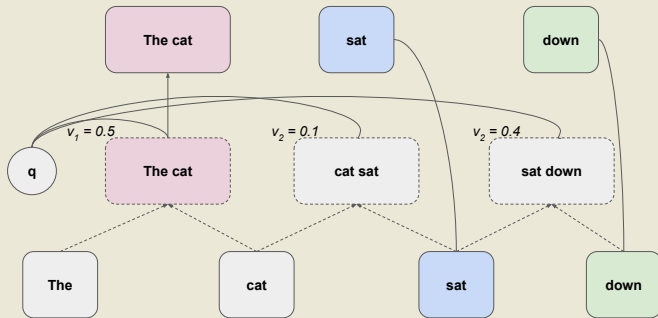
# Sampling from incremental structures

- Build a structure as a sequence of discrete choices (e.g., shift-reduce)
- Assigns a score to any (partial structure, action) tuple.
- Reparameterize the scores with Gumbel-Max - now we have a deterministic node.
- Forward: the **argmax** from the reparameterized scores for each step
- Backward: pretend we had used a **differentiable surrogate function**

Example: Gumbel Tree-LSTM [Choi et al., 2018].

# Example: Gumbel Tree-LSTM

- Building task-specific tree structures.
- Straight-Through Gumbel-Softmax at each step to select one arc.



# Sampling from factorized models

## Perturb-and-MAP

Reparameterize by **perturbing the arc scores**. (inexact!)

# Sampling from factorized models

## Perturb-and-MAP

Reparameterize by **perturbing the arc scores**. (inexact!)

- Sample from the standard Gumbel distribution.
- $\epsilon \sim G(0, 1)$

# Sampling from factorized models

## Perturb-and-MAP

Reparameterize by **perturbing the arc scores**. (inexact!)

- Sample from the standard Gumbel distribution.
- $\epsilon \sim G(0, 1)$
- Perturb the arc scores with the Gumbel noise.
- $\tilde{\eta} = \eta + \epsilon$

# Sampling from factorized models

## Perturb-and-MAP

Reparameterize by **perturbing the arc scores**. (inexact!)

- Sample from the standard Gumbel distribution.
- Perturb the arc scores with the Gumbel noise.
- Compute MAP (task-specific algorithm).
- $\epsilon \sim G(0, 1)$
- $\tilde{\eta} = \eta + \epsilon$
- $\arg \max_{\mathbf{z} \in \mathcal{Z}} \tilde{\eta}^T \mathbf{z}$

# Sampling from factorized models

## Perturb-and-MAP

Reparameterize by **perturbing the arc scores**. (inexact!)

- Sample from the standard Gumbel distribution.
- Perturb the arc scores with the Gumbel noise.
- Compute MAP (task-specific algorithm).
- Backward: we could use Straight-Through with Identity.
- $\epsilon \sim G(0, 1)$
- $\tilde{\eta} = \eta + \epsilon$
- $\arg \max_{\mathbf{z} \in \mathcal{Z}} \tilde{\eta}^T \mathbf{z}$

# Summary: Gradient surrogates

- Based on the **Straight-Through Estimator**.
- Can be used for stochastic or deterministic computation graphs.
- **Forward pass**: Get an argmax (might be structured).
- **Backpropagation**: use a function, which we hope is close to argmax.
- Examples:
  - Argmax for iterative structures and factorization into parts
  - Sampling from iterative structures and factorization into parts



# Gradient surrogates: Pros and cons

## Pros

- Do not suffer from the high variance problem of REINFORCE.
- Allow for flexibility to select or sample a latent structured in the middle of the computation graph.
- Efficient computation.

## Cons

- The Gumbel sampling with Perturb-and-MAP is an approximation.
- Bias, due to function mismatch on the backpropagation  
(next section will address this problem.)

# Overview

$$\mathbb{E}_{\pi_{\boldsymbol{\theta}}(\mathbf{z}|x)}[L(\mathbf{z})]$$

$$L(\arg \max_{\mathbf{z}} \pi_{\boldsymbol{\theta}}(\mathbf{z} | x))$$

- REINFORCE
- Straight-Through Gumbel (Perturb & MAP)
- Straight-Through
- SPIGOT

# Overview

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

$$L(\arg \max_{\mathbf{z}} \pi_{\theta}(\mathbf{z} | x))$$

$$L(\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[\mathbf{z}])$$

- REINFORCE
- Straight-Through Gumbel (Perturb & MAP)
- Straight-Through
- SPIGOT
- Structured Attn. Nets
- SparseMAP

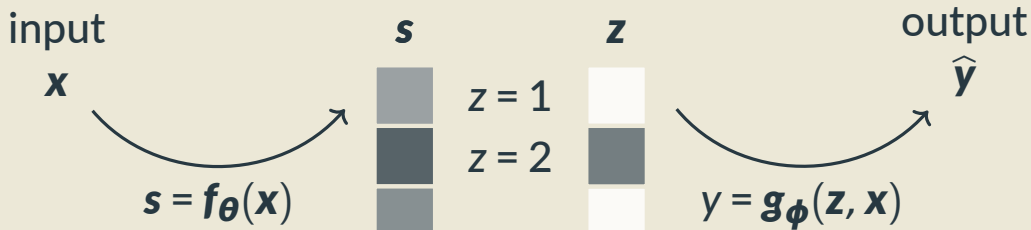
And more, in the next section!

# **IV. End-to-end Differentiable Relaxations**

# End-to-end differentiable relaxations

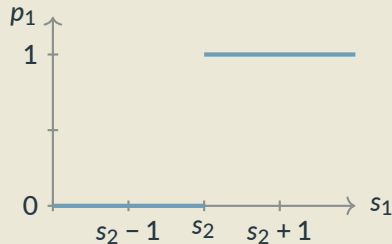
1. Digging into softmax
2. Alternatives to softmax
3. Generalizing to structured prediction
4. Stochasticity and global structures

# Recall: Discrete choices & differentiability

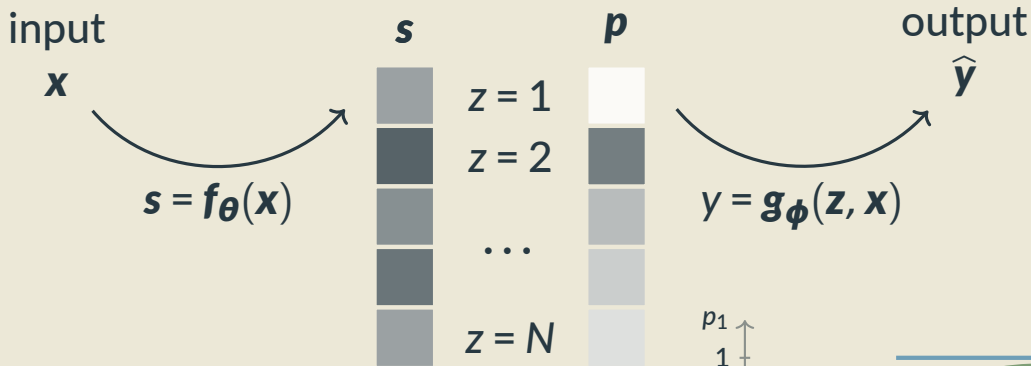


$$\frac{\partial \mathbf{z}}{\partial \mathbf{s}} = \mathbf{0} \text{ or n/a}$$

(argmax)




# One solution: smooth relaxation



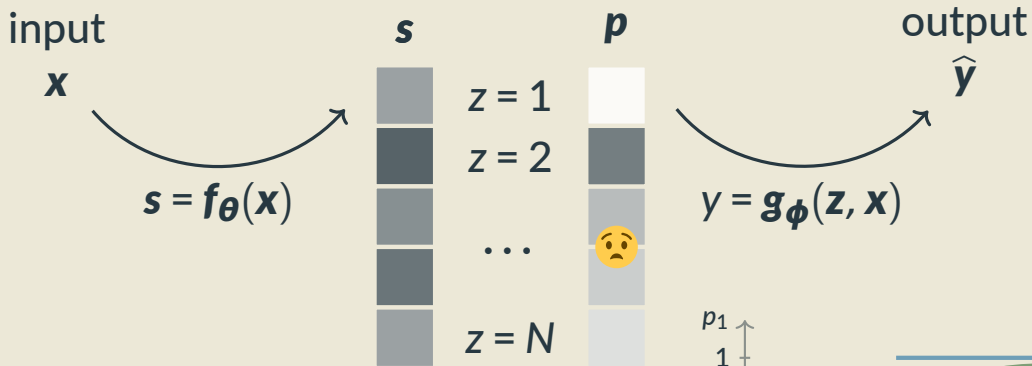
$\mathbf{p} = \text{softmax}(\mathbf{s}) = \mathbb{E}[\mathbf{z}]$ , i.e.  
replace  $\mathbb{E}[f(\mathbf{z})]$  with  $f(\mathbb{E}[\mathbf{z}])$

$$\frac{\partial \mathbf{p}}{\partial \mathbf{s}} = \text{😊}$$

(softmax)



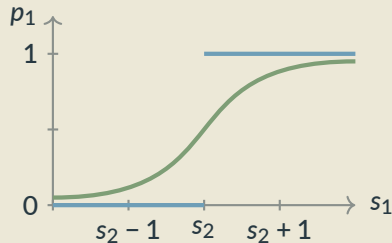
# One solution: smooth relaxation



$p = \text{softmax}(s) = \mathbb{E}[z]$ , i.e.  
replace  $\mathbb{E}[f(z)]$  with  $f(\mathbb{E}[z])$

$$\frac{\partial p}{\partial s} = \text{😊}$$

(softmax)





# Overview

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

$$L(\arg \max_{\mathbf{z}} \pi_{\theta}(\mathbf{z} | x))$$

$$L(\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[\mathbf{z}])$$

- REINFORCE
- Straight-Through Gumbel (Perturb & MAP)
- Straight-Through
- SPIGOT

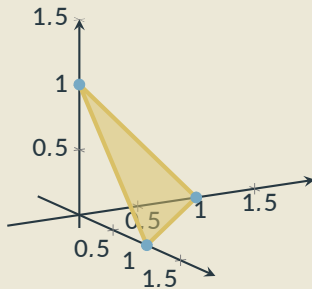
# What is softmax?

Often defined via  $p_i = \frac{\exp s_i}{\sum_j \exp s_j}$ , but where does it come from?

# What is softmax?

Often defined via  $p_i = \frac{\exp s_i}{\sum_j \exp s_j}$ , but where does it come from?

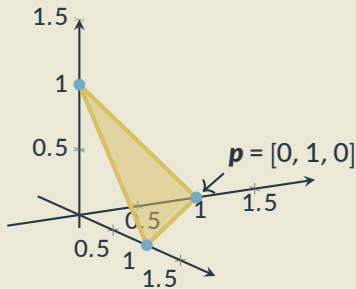
$\mathbf{p} \in \Delta$ : probability distribution over choices



# What is softmax?

Often defined via  $p_i = \frac{\exp s_i}{\sum_j \exp s_j}$ , but where does it come from?

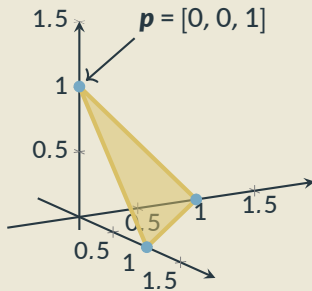
$\mathbf{p} \in \Delta$ : probability distribution over choices



# What is softmax?

Often defined via  $p_i = \frac{\exp s_i}{\sum_j \exp s_j}$ , but where does it come from?

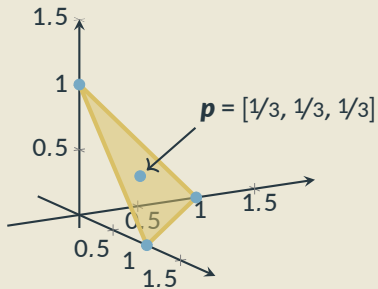
$\mathbf{p} \in \Delta$ : probability distribution over choices



# What is softmax?

Often defined via  $p_i = \frac{\exp s_i}{\sum_j \exp s_j}$ , but where does it come from?

$\mathbf{p} \in \Delta$ : probability distribution over choices

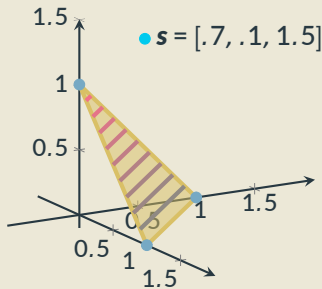


# What is softmax?

Often defined via  $p_i = \frac{\exp s_i}{\sum_j \exp s_j}$ , but where does it come from?

$\mathbf{p} \in \Delta$ : probability distribution over choices

Expected score under  $\mathbf{p}$ :  $\mathbb{E}_{i \sim \mathbf{p}} s_i = \mathbf{p}^\top \mathbf{s}$



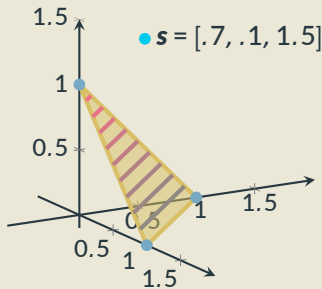
# What is softmax?

Often defined via  $p_i = \frac{\exp s_i}{\sum_j \exp s_j}$ , but where does it come from?

$\mathbf{p} \in \Delta$ : probability distribution over choices

Expected score under  $\mathbf{p}$ :  $\mathbb{E}_{i \sim \mathbf{p}} s_i = \mathbf{p}^\top \mathbf{s}$

**argmax**






# What is softmax?

Often defined via  $p_i = \frac{\exp s_i}{\sum_j \exp s_j}$ , but where does it come from?

$\mathbf{p} \in \Delta$ : probability distribution over choices

Expected score under  $\mathbf{p}$ :  $\mathbb{E}_{i \sim \mathbf{p}} s_i = \mathbf{p}^\top \mathbf{s}$

**argmax** maximizes **expected score**



# What is softmax?

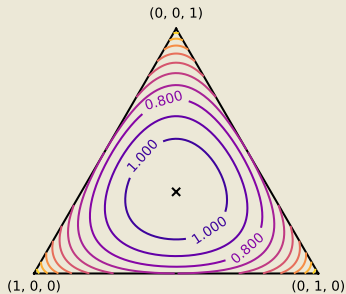
Often defined via  $p_i = \frac{\exp s_i}{\sum_j \exp s_j}$ , but where does it come from?

$\mathbf{p} \in \Delta$ : probability distribution over choices

Expected score under  $\mathbf{p}$ :  $\mathbb{E}_{i \sim \mathbf{p}} s_i = \mathbf{p}^\top \mathbf{s}$

**argmax** maximizes **expected score**

Shannon entropy of  $\mathbf{p}$ :  $H(\mathbf{p}) = -\sum_i p_i \log p_i$



# What is softmax?

Often defined via  $p_i = \frac{\exp s_i}{\sum_j \exp s_j}$ , but where does it come from?

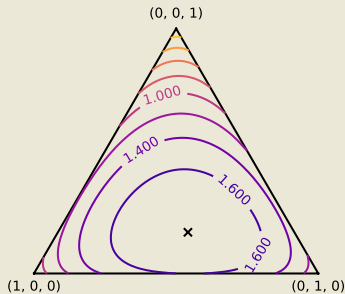
$\mathbf{p} \in \Delta$ : probability distribution over choices

Expected score under  $\mathbf{p}$ :  $\mathbb{E}_{i \sim \mathbf{p}} s_i = \mathbf{p}^\top \mathbf{s}$

**argmax** maximizes **expected score**

Shannon entropy of  $\mathbf{p}$ :  $H(\mathbf{p}) = -\sum_i p_i \log p_i$

**softmax** maximizes **expected score + entropy**:



$$\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$$

# Variational form of softmax

**Proposition.** The unique solution to  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$  is given by  $p_j = \frac{\exp s_j}{\sum_i \exp s_i}$ .

# Variational form of softmax

**Proposition.** The unique solution to  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$  is given by  $p_j = \frac{\exp s_j}{\sum_i \exp s_i}$ .

Explicit form of the optimization problem:

$$\begin{aligned} &\text{maximize} && \sum_j p_j s_j - p_j \log p_j \\ &\text{subject to} && \mathbf{p} \geq 0, \mathbf{p}^\top \mathbf{1} = 1 \end{aligned}$$

# Variational form of softmax

**Proposition.** The unique solution to  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$  is given by  $p_j = \frac{\exp s_j}{\sum_i \exp s_i}$ .

Explicit form of the optimization problem:

$$\begin{aligned} &\text{maximize} && \sum_j p_j s_j - p_j \log p_j \\ &\text{subject to} && \mathbf{p} \geq 0, \mathbf{p}^\top \mathbf{1} = 1 \end{aligned}$$

Lagrangian:

$$\mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = - \sum_j p_j s_j - p_j \log p_j - \mathbf{p}^\top \boldsymbol{\nu} + \tau(\mathbf{p}^\top \mathbf{1} - 1)$$

# Variational form of softmax

**Proposition.** The unique solution to  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$  is given by  $p_j = \frac{\exp s_j}{\sum_i \exp s_i}$ .

Explicit form of the optimization problem:

$$\begin{aligned} & \text{maximize} && \sum_j p_j s_j - p_j \log p_j \\ & \text{subject to} && \mathbf{p} \geq 0, \mathbf{p}^\top \mathbf{1} = 1 \end{aligned}$$

Lagrangian:

$$\mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -\sum_j p_j s_j - p_j \log p_j - \mathbf{p}^\top \boldsymbol{\nu} + \tau(\mathbf{p}^\top \mathbf{1} - 1)$$

Optimality conditions (KKT):

$$0 = \nabla_{p_i} \mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -s_i + \log p_i + 1 - \nu_i + \tau$$

$$\mathbf{p}^\top \boldsymbol{\nu} = 0$$

$$\mathbf{p} \in \Delta$$

$$\boldsymbol{\nu} \geq 0$$

# Variational form of softmax

**Proposition.** The unique solution to  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$  is given by  $p_j = \frac{\exp s_j}{\sum_i \exp s_i}$ .

Explicit form of the optimization problem:

$$\log p_i = s_i + \nu_i - (\tau + 1)$$

$$\text{maximize} \quad \sum_j p_j s_j - p_j \log p_j$$

$$\text{subject to} \quad \mathbf{p} \geq 0, \mathbf{p}^\top \mathbf{1} = 1$$

Lagrangian:

$$\mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -\sum_j p_j s_j - p_j \log p_j - \mathbf{p}^\top \boldsymbol{\nu} + \tau(\mathbf{p}^\top \mathbf{1} - 1)$$

Optimality conditions (KKT):

$$0 = \nabla_{p_i} \mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -s_i + \log p_i + 1 - \nu_i + \tau$$

$$\mathbf{p}^\top \boldsymbol{\nu} = 0$$

$$\mathbf{p} \in \Delta$$

$$\boldsymbol{\nu} \geq 0$$



# Variational form of softmax

**Proposition.** The unique solution to  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$  is given by  $p_j = \frac{\exp s_j}{\sum_i \exp s_i}$ .

Explicit form of the optimization problem:

$$\begin{aligned} & \text{maximize} && \sum_j p_j s_j - p_j \log p_j \\ & \text{subject to} && \mathbf{p} \geq 0, \mathbf{p}^\top \mathbf{1} = 1 \end{aligned}$$

$$\begin{aligned} \log p_i &= s_i + \nu_i - (\tau + 1) \\ \text{if } p_i &= 0, \text{ r.h.s. must be } -\infty, \\ \text{thus } p_i &> 0, \text{ so } \nu_i = 0. \end{aligned}$$

Lagrangian:

$$\mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -\sum_j p_j s_j - p_j \log p_j - \mathbf{p}^\top \boldsymbol{\nu} + \tau(\mathbf{p}^\top \mathbf{1} - 1)$$

Optimality conditions (KKT):

$$0 = \nabla_{p_i} \mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -s_i + \log p_i + 1 - \nu_i + \tau$$

$$\mathbf{p}^\top \boldsymbol{\nu} = 0$$

$$\mathbf{p} \in \Delta$$

$$\boldsymbol{\nu} \geq 0$$

# Variational form of softmax

**Proposition.** The unique solution to  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$  is given by  $p_j = \frac{\exp s_j}{\sum_i \exp s_i}$ .

Explicit form of the optimization problem:

$$\begin{aligned} & \text{maximize} && \sum_j p_j s_j - p_j \log p_j \\ & \text{subject to} && \mathbf{p} \geq 0, \mathbf{p}^\top \mathbf{1} = 1 \end{aligned}$$

$$\begin{aligned} \log p_i &= s_i + \nu_i - (\tau + 1) \\ \text{if } p_i &= 0, \text{ r.h.s. must be } -\infty, \\ \text{thus } p_i &> 0, \text{ so } \nu_i = 0. \end{aligned}$$

$$p_i = \exp(s_i) / \exp(\tau + 1) = \exp(s_i) / Z$$

Lagrangian:

$$\mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -\sum_j p_j s_j - p_j \log p_j - \mathbf{p}^\top \boldsymbol{\nu} + \tau(\mathbf{p}^\top \mathbf{1} - 1)$$

Optimality conditions (KKT):

$$0 = \nabla_{p_i} \mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -s_i + \log p_i + 1 - \nu_i + \tau$$

$$\mathbf{p}^\top \boldsymbol{\nu} = 0$$

$$\mathbf{p} \in \Delta$$

$$\boldsymbol{\nu} \geq 0$$

# Variational form of softmax

**Proposition.** The unique solution to  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$  is given by  $p_j = \frac{\exp s_j}{\sum_i \exp s_i}$ .

Explicit form of the optimization problem:

$$\begin{aligned} & \text{maximize} && \sum_j p_j s_j - p_j \log p_j \\ & \text{subject to} && \mathbf{p} \geq 0, \mathbf{p}^\top \mathbf{1} = 1 \end{aligned}$$

Lagrangian:

$$\mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -\sum_j p_j s_j - p_j \log p_j - \mathbf{p}^\top \boldsymbol{\nu} + \tau(\mathbf{p}^\top \mathbf{1} - 1)$$

Optimality conditions (KKT):

$$\begin{aligned} 0 &= \nabla_{p_i} \mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -s_i + \log p_i + 1 - \nu_i + \tau \\ \mathbf{p}^\top \boldsymbol{\nu} &= 0 \\ \mathbf{p} &\in \Delta \\ \boldsymbol{\nu} &\geq 0 \end{aligned}$$

$$\begin{aligned} \log p_i &= s_i + \nu_i - (\tau + 1) \\ \text{if } p_i &= 0, \text{ r.h.s. must be } -\infty, \\ \text{thus } p_i &> 0, \text{ so } \nu_i = 0. \end{aligned}$$

$$p_i = \exp(s_i) / \exp(\tau + 1) = \exp(s_i) / Z$$

Must find  $Z$  such that  $\sum_j p_j = 1$ .

# Variational form of softmax

**Proposition.** The unique solution to  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$  is given by  $p_j = \frac{\exp s_j}{\sum_i \exp s_i}$ .

Explicit form of the optimization problem:

$$\begin{aligned} & \text{maximize} && \sum_j p_j s_j - p_j \log p_j \\ & \text{subject to} && \mathbf{p} \geq 0, \mathbf{p}^\top \mathbf{1} = 1 \end{aligned}$$

Lagrangian:

$$\mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -\sum_j p_j s_j - p_j \log p_j - \mathbf{p}^\top \boldsymbol{\nu} + \tau(\mathbf{p}^\top \mathbf{1} - 1)$$

Optimality conditions (KKT):

$$\begin{aligned} 0 &= \nabla_{p_i} \mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -s_i + \log p_i + 1 - \nu_i + \tau \\ \mathbf{p}^\top \boldsymbol{\nu} &= 0 \\ \mathbf{p} &\in \Delta \\ \boldsymbol{\nu} &\geq 0 \end{aligned}$$

$$\begin{aligned} \log p_i &= s_i + \nu_i - (\tau + 1) \\ \text{if } p_i &= 0, \text{ r.h.s. must be } -\infty, \\ \text{thus } p_i &> 0, \text{ so } \nu_i = 0. \end{aligned}$$

$$p_i = \exp(s_i) / \exp(\tau + 1) = \exp(s_i) / Z$$

Must find  $Z$  such that  $\sum_j p_j = 1$ .

Answer:  $Z = \sum_j \exp(s_j)$

# Variational form of softmax

**Proposition.** The unique solution to  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^\top \mathbf{s} + H(\mathbf{p})$  is given by  $p_j = \frac{\exp s_j}{\sum_i \exp s_i}$ .

Explicit form of the optimization problem:

$$\begin{aligned} & \text{maximize} && \sum_j p_j s_j - p_j \log p_j \\ & \text{subject to} && \mathbf{p} \geq 0, \mathbf{p}^\top \mathbf{1} = 1 \end{aligned}$$

Lagrangian:

$$\mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -\sum_j p_j s_j - p_j \log p_j - \mathbf{p}^\top \boldsymbol{\nu} + \tau(\mathbf{p}^\top \mathbf{1} - 1)$$

Optimality conditions (KKT):

$$0 = \nabla_{p_i} \mathcal{L}(\mathbf{p}, \boldsymbol{\nu}, \tau) = -s_i + \log p_i + 1 - \nu_i + \tau$$

$$\mathbf{p}^\top \boldsymbol{\nu} = 0$$

$$\mathbf{p} \in \Delta$$

$$\boldsymbol{\nu} \geq 0$$

$$\begin{aligned} \log p_i &= s_i + \nu_i - (\tau + 1) \\ \text{if } p_i &= 0, \text{ r.h.s. must be } -\infty, \\ \text{thus } p_i &> 0, \text{ so } \nu_i = 0. \end{aligned}$$

$$p_i = \exp(s_i) / \exp(\tau + 1) = \exp(s_i) / Z$$

Must find  $Z$  such that  $\sum_j p_j = 1$ .

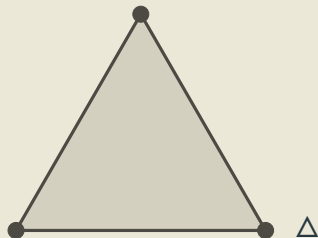

Answer:  $Z = \sum_j \exp(s_j)$

$$\text{So, } p_i = \frac{\exp(s_i)}{\sum_j \exp(s_j)}.$$

Classic result, e.g., [Boyd and Vandenberghe, 2004, Wainwright and Jordan, 2008]

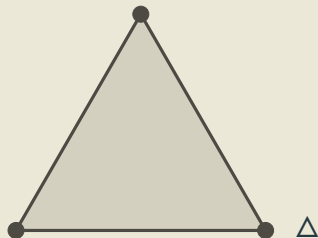

# Generalizing softmax: Smoothed argmaxes

$$\hat{\mathbf{p}}_{\Omega}(\mathbf{s}) = \arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^{\top} \mathbf{s} - \Omega(\mathbf{p})$$



# Generalizing softmax: Smoothed argmaxes

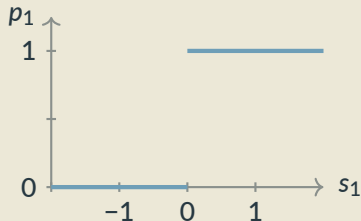
$$\hat{\mathbf{p}}_{\Omega}(\mathbf{s}) = \arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^{\top} \mathbf{s} - \Omega(\mathbf{p})$$




# Generalizing softmax: Smoothed argmaxes

$$\hat{\mathbf{p}}_{\Omega}(\mathbf{s}) = \arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^{\top} \mathbf{s} - \Omega(\mathbf{p})$$

- argmax:  $\Omega(\mathbf{p}) = 0$



$[0, 0, 1]$


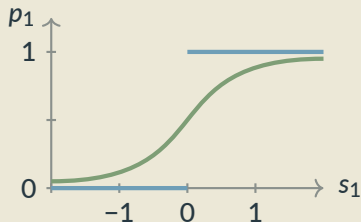




# Generalizing softmax: Smoothed argmaxes

$$\hat{\mathbf{p}}_{\Omega}(\mathbf{s}) = \arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^{\top} \mathbf{s} - \Omega(\mathbf{p})$$

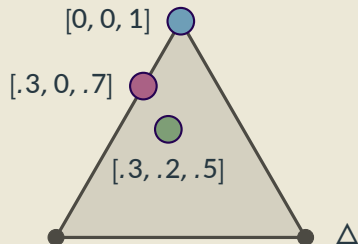

- argmax:  $\Omega(\mathbf{p}) = 0$
- softmax:  $\Omega(\mathbf{p}) = \sum_j p_j \log p_j$



# Generalizing softmax: Smoothed argmaxes

$$\hat{\mathbf{p}}_{\Omega}(\mathbf{s}) = \arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^{\top} \mathbf{s} - \Omega(\mathbf{p})$$

- argmax:  $\Omega(\mathbf{p}) = 0$
- softmax:  $\Omega(\mathbf{p}) = \sum_j p_j \log p_j$
- sparsemax:  $\Omega(\mathbf{p}) = 1/2 \|\mathbf{p}\|_2^2$




# Generalizing softmax: Smoothed argmaxes

$$\hat{\mathbf{p}}_{\Omega}(\mathbf{s}) = \arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^{\top} \mathbf{s} - \Omega(\mathbf{p})$$

- argmax:  $\Omega(\mathbf{p}) = 0$
- softmax:  $\Omega(\mathbf{p}) = \sum_j p_j \log p_j$
- sparsemax:  $\Omega(\mathbf{p}) = 1/2 \|\mathbf{p}\|_2^2$
- $\alpha$ -entmax:  $\Omega(\mathbf{p}) = 1/\alpha(\alpha-1) \sum_j p_j^{\alpha}$


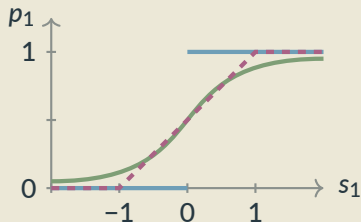
Generalized entropy interpolates in between [Tsallis, 1988]  
 Used in Sparse Seq2Seq: [Peters et al., 2019] and Adaptively  
 Sparse Transformers [Correia et al., 2019]



# Generalizing softmax: Smoothed argmaxes

$$\hat{\mathbf{p}}_{\Omega}(\mathbf{s}) = \arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^{\top} \mathbf{s} - \Omega(\mathbf{p})$$

- argmax:  $\Omega(\mathbf{p}) = 0$
- softmax:  $\Omega(\mathbf{p}) = \sum_j p_j \log p_j$
- sparsemax:  $\Omega(\mathbf{p}) = 1/2 \|\mathbf{p}\|_2^2$
- $\alpha$ -entmax:  $\Omega(\mathbf{p}) = 1/\alpha(\alpha-1) \sum_j p_j^{\alpha}$
- fusedmax:  $\Omega(\mathbf{p}) = 1/2 \|\mathbf{p}\|_2^2 + \sum_j |p_j - p_{j-1}|$
- csparsesmax:  $\Omega(\mathbf{p}) = 1/2 \|\mathbf{p}\|_2^2 + \iota(\mathbf{a} \leq \mathbf{p} \leq \mathbf{b})$
- csoftmax:  $\Omega(\mathbf{p}) = \sum_j p_j \log p_j + \iota(\mathbf{a} \leq \mathbf{p} \leq \mathbf{b})$



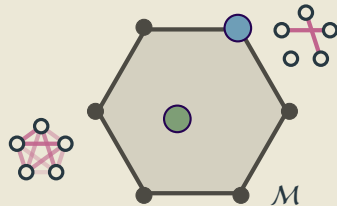
# The structured case: Marginal polytope


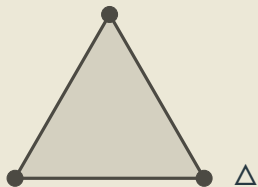
- Each vertex corresponds to one such *bit vector*  $\mathbf{z}$
- Points inside correspond to *marginal distributions*: convex combinations of structured objects

$$\boldsymbol{\mu} = \underbrace{p_1 \mathbf{z}_1 + \dots + p_N \mathbf{z}_N}_{\text{exponentially many terms}}, \quad \mathbf{p} \in \Delta.$$

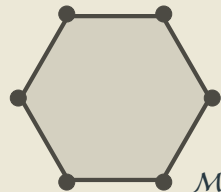
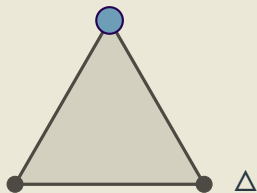
$$\begin{aligned} p_1 &= 0.2, & \mathbf{z}_1 &= [1, 0, 0, 0, 1, 0, 0, 0, 1] \\ p_2 &= 0.7, & \mathbf{z}_2 &= [0, 0, 1, 0, 0, 1, 1, 0, 0] \\ p_3 &= 0.1, & \mathbf{z}_3 &= [1, 0, 0, 0, 1, 0, 0, 1, 0] \end{aligned}$$

$$\Rightarrow \boldsymbol{\mu} = [.3, 0, .7, 0, .3, .7, .7, .1, .2].$$

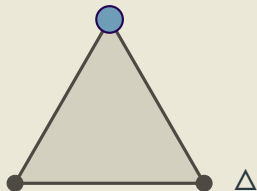




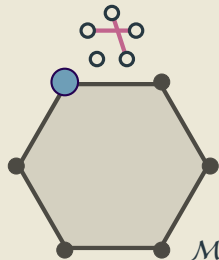
- $\operatorname{argmax}_{\mathbf{p} \in \Delta} \arg \max \mathbf{p}^T \mathbf{s}$



•  $\text{argmax}_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s}$



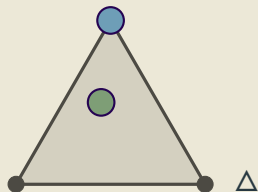
•  $\text{MAP}_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta}$



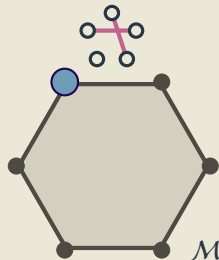


- **argmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s}$

- **softmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s} + H(\mathbf{p})$

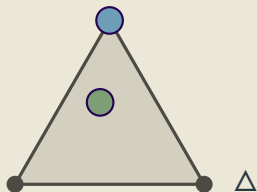


- **MAP**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta}$




- **argmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s}$

- **softmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s} + H(\mathbf{p})$



- **MAP**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta}$

- **marginals**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta} + \tilde{H}(\boldsymbol{\mu})$



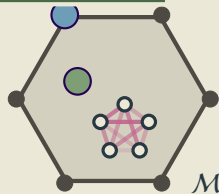
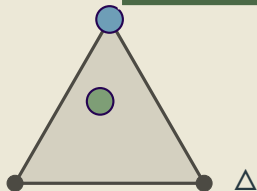
- **argmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s}$

- **softmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s} + H(\mathbf{p})$

- **MAP**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta}$

- **marginals**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta} + \tilde{H}(\boldsymbol{\mu})$

Just like softmax relaxes argmax,  
marginals relax MAP **differentiably!**



- **argmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s}$

- **softmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s} + H(\mathbf{p})$

- **MAP**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta}$

- **marginals**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta} + \tilde{H}(\boldsymbol{\mu})$

Just like softmax relaxes argmax,  
marginals relax MAP **differentiably!**

Unlike argmax/softmax, computation is not obvious!



# Algorithms for specific structures

	Best structure (MAP)	Marginals
Sequence tagging	Viterbi [Rabiner, 1989]	Forward-Backward [Rabiner, 1989]
Constituent trees	CKY [Kasami, 1966, Younger, 1967] [Cocke and Schwartz, 1970]	Inside-Outside [Baker, 1979]
Temporal alignments	DTW [Sakoe and Chiba, 1978]	Soft-DTW [Cuturi and Blondel, 2017]
Dependency trees	Max. Spanning Arborescence [Chu and Liu, 1965, Edmonds, 1967]	Matrix-Tree [Kirchhoff, 1847]
Assignments	Kuhn-Munkres [Kuhn, 1955, Jonker and Volgenant, 1987]	#P-complete [Valiant, 1979, Taskar, 2004]

# Algorithms for specific structures

		Best structure (MAP)	Marginals
dyn. prog.	Sequence tagging	Viterbi [Rabiner, 1989]	Forward-Backward [Rabiner, 1989]
	Constituent trees	CKY [Kasami, 1966, Younger, 1967] [Cocke and Schwartz, 1970]	Inside-Outside [Baker, 1979]
	Temporal alignments	DTW [Sakoe and Chiba, 1978]	Soft-DTW [Cuturi and Blondel, 2017]
	Dependency trees	Max. Spanning Arborescence [Chu and Liu, 1965, Edmonds, 1967]	Matrix-Tree [Kirchhoff, 1847]
	Assignments	Kuhn-Munkres [Kuhn, 1955, Jonker and Volgenant, 1987]	#P-complete [Valiant, 1979, Taskar, 2004]

# Derivatives of marginals 1: DP

**Dynamic programming:** marginals by **Forward-Backward**, **Inside-Outside**, etc.

# Derivatives of marginals 1: DP

**Dynamic programming:** marginals by **Forward-Backward, Inside-Outside**, etc.

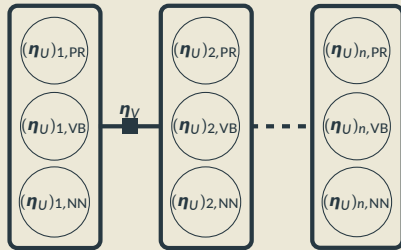
---

## Marginals in a sequence tagging model.

---

```
1 input:  $d$  tags,  $n$  tokens,  $\boldsymbol{\eta}_U \in \mathbb{R}^{n \times d}$ ,  $\boldsymbol{\eta}_V \in \mathbb{R}^{d \times d}$ 
2 initialize  $\boldsymbol{\alpha}_1 = \mathbf{0}$ ,  $\boldsymbol{\beta}_n = \mathbf{0}$ 
3 for  $i \in 2, \dots, n$  do                                # forward log-probabilities
4    $\alpha_{i,k} = \log \sum_{k'} \exp(\alpha_{i-1,k'} + (\boldsymbol{\eta}_U)_{i,k} + (\boldsymbol{\eta}_V)_{k',k})$  for all  $k$ 
5 for  $i \in n-1, \dots, 1$  do                                # backward log-probabilities
6    $\beta_{i,k} = \log \sum_{k'} \exp(\beta_{i+1,k'} + (\boldsymbol{\eta}_U)_{i+1,k'} + (\boldsymbol{\eta}_V)_{k,k'})$  for all  $k$ 
7  $Z = \sum_k \exp \alpha_{n,k}$                                     # partition function
8 return  $\boldsymbol{\mu} = \exp(\boldsymbol{\alpha} + \boldsymbol{\beta} - \log Z)$                 # marginals
```

---





# Derivatives of marginals 1: DP

**Dynamic programming:** marginals by **Forward-Backward, Inside-Outside**, etc.

- Alg. consists of differentiable ops: PyTorch autograd can handle it! (v. bad idea)

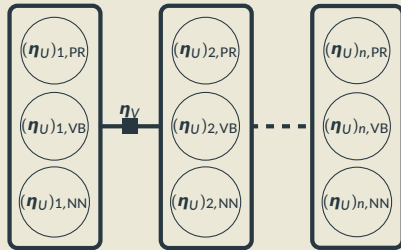
---

## Marginals in a sequence tagging model.

---

```
1 input:  $d$  tags,  $n$  tokens,  $\boldsymbol{\eta}_U \in \mathbb{R}^{n \times d}$ ,  $\boldsymbol{\eta}_V \in \mathbb{R}^{d \times d}$ 
2 initialize  $\boldsymbol{\alpha}_1 = \mathbf{0}$ ,  $\boldsymbol{\beta}_n = \mathbf{0}$ 
3 for  $i \in 2, \dots, n$  do                                # forward log-probabilities
4    $\alpha_{i,k} = \log \sum_{k'} \exp(\alpha_{i-1,k'} + (\boldsymbol{\eta}_U)_{i,k} + (\boldsymbol{\eta}_V)_{k',k})$    for all  $k$ 
5 for  $i \in n-1, \dots, 1$  do                                # backward log-probabilities
6    $\beta_{i,k} = \log \sum_{k'} \exp(\beta_{i+1,k'} + (\boldsymbol{\eta}_U)_{i+1,k'} + (\boldsymbol{\eta}_V)_{k,k'})$    for all  $k$ 
7  $Z = \sum_k \exp \alpha_{n,k}$                                 # partition function
8 return  $\boldsymbol{\mu} = \exp(\boldsymbol{\alpha} + \boldsymbol{\beta} - \log Z)$                 # marginals
```

---



# Derivatives of marginals 1: DP

**Dynamic programming:** marginals by **Forward-Backward, Inside-Outside**, etc.

- Alg. consists of differentiable ops: PyTorch autograd can handle it! (v. bad idea)
- Better book-keeping: Li and Eisner [2009], Mensch and Blondel [2018]


---

## Marginals in a sequence tagging model.

---

```
1 input:  $d$  tags,  $n$  tokens,  $\boldsymbol{\eta}_U \in \mathbb{R}^{n \times d}$ ,  $\boldsymbol{\eta}_V \in \mathbb{R}^{d \times d}$ 
2 initialize  $\boldsymbol{\alpha}_1 = \mathbf{0}$ ,  $\boldsymbol{\beta}_n = \mathbf{0}$ 
3 for  $i \in 2, \dots, n$  do                                # forward log-probabilities
4    $\alpha_{i,k} = \log \sum_{k'} \exp(\alpha_{i-1,k'} + (\boldsymbol{\eta}_U)_{i,k} + (\boldsymbol{\eta}_V)_{k',k})$    for all  $k$ 
5 for  $i \in n-1, \dots, 1$  do                                # backward log-probabilities
6    $\beta_{i,k} = \log \sum_{k'} \exp(\beta_{i+1,k'} + (\boldsymbol{\eta}_U)_{i+1,k'} + (\boldsymbol{\eta}_V)_{k,k'})$    for all  $k$ 
7  $Z = \sum_k \exp \alpha_{n,k}$                                 # partition function
8 return  $\boldsymbol{\mu} = \exp(\boldsymbol{\alpha} + \boldsymbol{\beta} - \log Z)$                 # marginals
```

---



# Derivatives of marginals 1: DP

**Dynamic programming:** marginals by **Forward-Backward, Inside-Outside**, etc.

- Alg. consists of differentiable ops: PyTorch autograd can handle it! (v. bad idea)
- Better book-keeping: Li and Eisner [2009], Mensch and Blondel [2018]
- With circular dependencies, this breaks! Can get an approximation [Stoyanov et al., 2011]

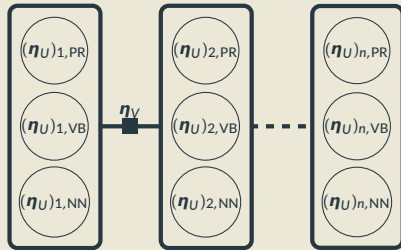
---

Marginals in a sequence tagging model.

---

```
1 input:  $d$  tags,  $n$  tokens,  $\boldsymbol{\eta}_U \in \mathbb{R}^{n \times d}$ ,  $\boldsymbol{\eta}_V \in \mathbb{R}^{d \times d}$ 
2 initialize  $\boldsymbol{\alpha}_1 = \mathbf{0}$ ,  $\boldsymbol{\beta}_n = \mathbf{0}$ 
3 for  $i \in 2, \dots, n$  do                                # forward log-probabilities
4    $\alpha_{i,k} = \log \sum_{k'} \exp(\alpha_{i-1,k'} + (\boldsymbol{\eta}_U)_{i,k} + (\boldsymbol{\eta}_V)_{k',k})$    for all  $k$ 
5 for  $i \in n-1, \dots, 1$  do                                # backward log-probabilities
6    $\beta_{i,k} = \log \sum_{k'} \exp(\beta_{i+1,k'} + (\boldsymbol{\eta}_U)_{i+1,k'} + (\boldsymbol{\eta}_V)_{k,k'})$    for all  $k$ 
7  $Z = \sum_k \exp \alpha_{n,k}$                                 # partition function
8 return  $\boldsymbol{\mu} = \exp(\boldsymbol{\alpha} + \boldsymbol{\beta} - \log Z)$                 # marginals
```

---



# Derivatives of marginals 2: Matrix-Tree

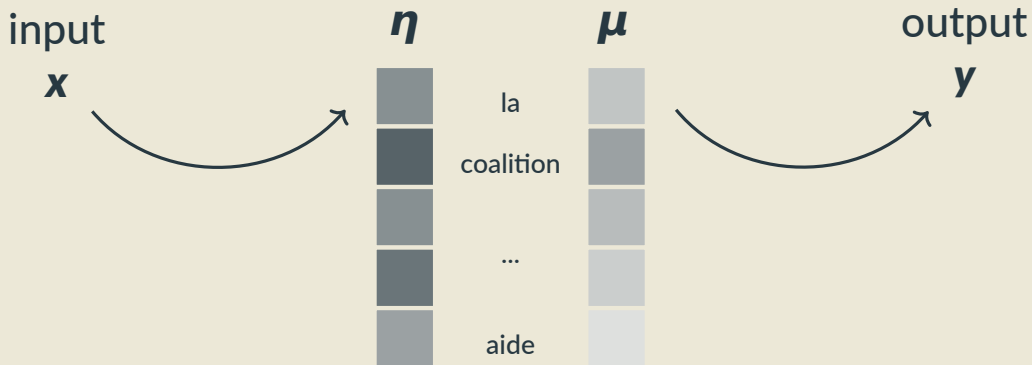
$\mathbf{L}(\mathbf{s})$ : Laplacian of the edge score graph

$$Z = \det \mathbf{L}(\mathbf{s})$$

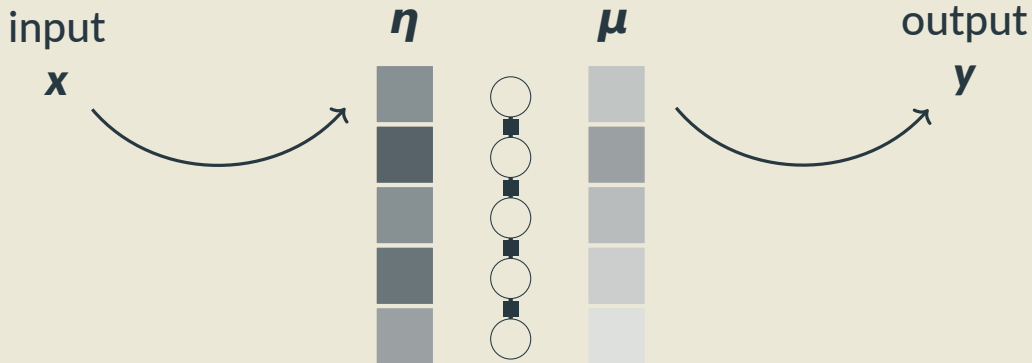
$$\boldsymbol{\mu} = \mathbf{L}(\mathbf{s})^{-1}$$

$$\nabla \boldsymbol{\mu} = \nabla \mathbf{L}^{-1} = \mathbf{L}^{-1} \left( \frac{\partial \mathbf{L}}{\partial \boldsymbol{\eta}} \right) \mathbf{L}^{-1}$$

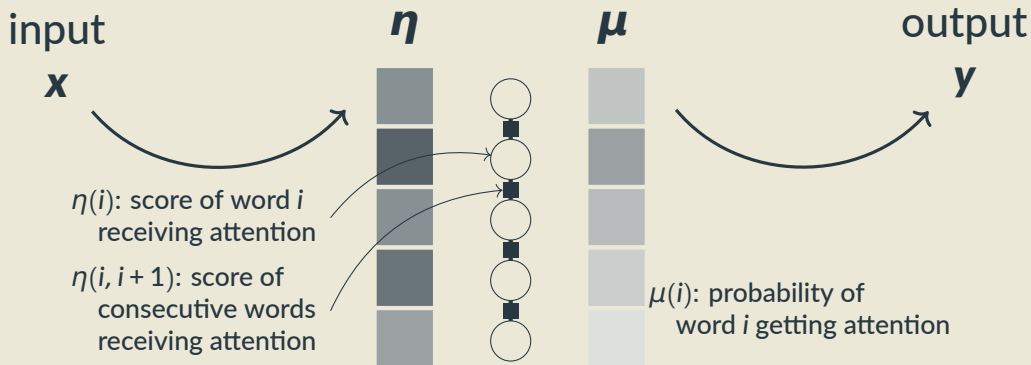
# Structured Attention Networks



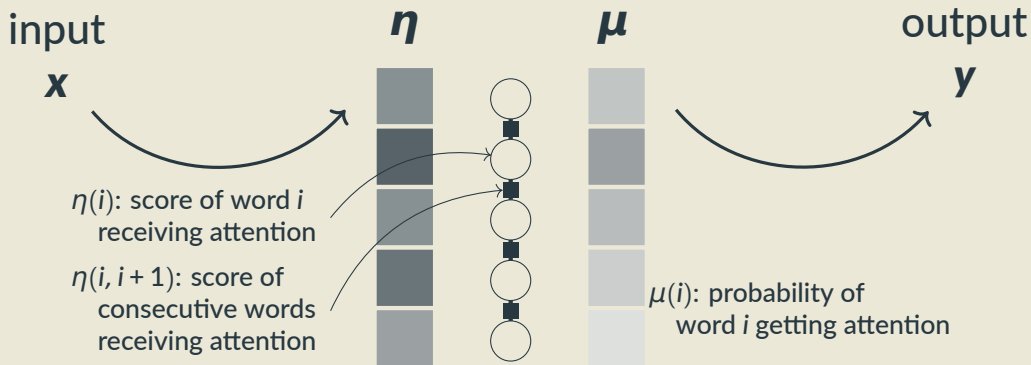
# Structured Attention Networks



# Structured Attention Networks



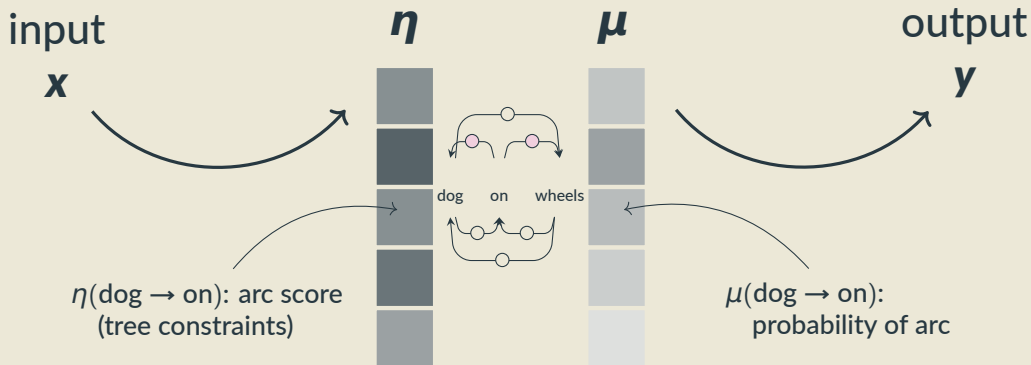
# Structured Attention Networks



CRF marginals (from *forward-backward*) give attention weights  $\in (0, 1)$




# Structured Attention Networks



CRF marginals (from *forward-backward*) give attention weights  $\in (0, 1)$   
 Similar idea for projective dependency trees with *inside-outside*

# Structured Attention Networks



CRF marginals (from *forward-backward*) give attention weights  $\in (0, 1)$

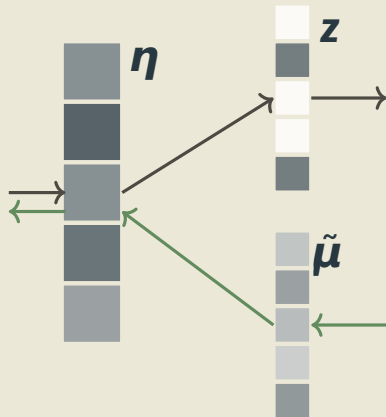
Similar idea for projective dependency trees with *inside-outside*

and non-projective with the Matrix-Tree theorem [Liu and Lapata, 2018].

# Differentiable Perturb & Parse

## Extending Gumbel-Softmax to structured stochastic models

- Forward pass:  
sample structure  $\mathbf{z}$  (approximately)  
$$\mathbf{z} = \arg \max_{\mathbf{z} \in \mathcal{Z}} (\boldsymbol{\eta} + \boldsymbol{\epsilon})^\top \mathbf{z}$$
- Backward pass:  
pretend we did marginal inference  
$$\tilde{\boldsymbol{\mu}} = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} (\boldsymbol{\eta} + \boldsymbol{\epsilon})^\top \mathbf{z} + \tilde{H}(\boldsymbol{\mu})$$
  
(or some similar relaxation)



# Back-propagating through marginals

Pros:

# Back-propagating through marginals

Pros:

- Familiar algorithms for NLPers,

# Back-propagating through marginals

Pros:

- Familiar algorithms for NLPers,
- (Structured Attention Networks:) All computations exact.

# Back-propagating through marginals

Pros:

- Familiar algorithms for NLPers,
- (Structured Attention Networks:) All computations exact.

Cons:

- (Structured Attention Networks:) forward pass marginals are dense;  
(fixed by Perturb & MAP, at cost of rough approximation)

# Back-propagating through marginals

Pros:

- Familiar algorithms for NLPers,
- (Structured Attention Networks:) All computations exact.

Cons:

- (Structured Attention Networks:) forward pass marginals are dense;  
(fixed by Perturb & MAP, at cost of rough approximation)
- Efficient & numerically stable back-propagation through DPs is tricky;  
(somewhat alleviated by Mensch and Blondel [2018])



# Back-propagating through marginals

## Pros:

- Familiar algorithms for NLPers,
- (Structured Attention Networks:) All computations exact.

## Cons:

- (Structured Attention Networks:) forward pass marginals are dense;  
(fixed by Perturb & MAP, at cost of rough approximation)
- Efficient & numerically stable back-propagation through DPs is tricky;  
(somewhat alleviated by Mensch and Blondel [2018])
- Not applicable when marginals are unavailable.

# Back-propagating through marginals

## Pros:

- Familiar algorithms for NLPers,
- (Structured Attention Networks:) All computations exact.

## Cons:

- (Structured Attention Networks:) forward pass marginals are dense;  
(fixed by Perturb & MAP, at cost of rough approximation)
- Efficient & numerically stable back-propagation through DPs is tricky;  
(somewhat alleviated by Mensch and Blondel [2018])
- Not applicable when marginals are unavailable.
- Case-by-case algorithms required, can get tedious.

# Back-propagating through marginals

Pros:

- Familiar algorithms for NLPers,
- (Structured Attention Net

Cons:

- (Structured Attention Net (fixed by Perturb & MA
- Efficient & numerically stable (somewhat alleviated b
- Not applicable when mar
- Case-by-case algorithms

```

procedure BACKPROPINSIDEOUTSIDE( $\theta, p, \nabla_p^T$ )
  for  $s, t = 1, \dots, n, s \neq t$  do
     $d[s, t] \leftarrow \log p[s, t] \otimes \log \nabla_p^T[s, t]$ 
     $\nabla_p^T[s, t] \leftarrow \log \nabla_p^T[s, t] \leftarrow -\infty$ 
    for  $s = 1, \dots, n - 1$  do
      for  $t = s + 1, \dots, n$  do
         $\nabla_p^T[s, t, R, 0], \nabla_p^T[s, t, R, 1] \leftarrow d[s, t]$ 
         $\nabla_p^T[1, n, R, 1] \leftarrow -d[s, t]$ 
        if  $s > 1$  then
           $\nabla_p^T[s, t, L, 0], \nabla_p^T[s, t, L, 1] \leftarrow -d[s, t]$ 
           $\nabla_p^T[1, n, R, 1] \leftarrow -d[s, t]$ 
    for  $k = 1, \dots, n$  do
      for  $s = 1, \dots, n - k$  do
         $t \leftarrow s + k$ 
         $v \leftarrow \nabla_p^T[s, t, R, 0] \otimes \beta[p, t, R, 0]$ 
        for  $u = t, \dots, n$  do
           $\nabla_p^T[s, u, R, 0], \nabla_p^T[s, u, R, 1] \leftarrow v \otimes \beta[u, u, R, 1] \otimes \alpha[t, u, R, 1]$ 
        if  $s > 1$  then
           $v \leftarrow \nabla_p^T[s, t, L, 0] \otimes \beta[p, t, L, 0]$ 
          for  $u = t, \dots, n$  do
             $\nabla_p^T[s, u, L, 0], \nabla_p^T[s, u, L, 1] \leftarrow v \otimes \beta[u, u, L, 1] \otimes \alpha[t, u, L, 1]$ 
           $v \leftarrow \nabla_p^T[s, t, L, 1] \otimes \beta[p, t, L, 1]$ 
          for  $u = t, \dots, n$  do
             $\nabla_p^T[s, u, L, 1], \nabla_p^T[s, u, L, 0] \leftarrow v \otimes \beta[u, u, L, 0] \otimes \alpha[t, u, L, 0]$ 
          for  $u = 1, \dots, s - 1$  do
             $\gamma \leftarrow \beta[u, t, R, 0] \otimes \alpha[u, s - 1, R, 1] \otimes \theta_{s, t}$ 
             $\nabla_p^T[s, t, R, 0], \nabla_p^T[s, t, R, 1] \leftarrow \gamma \otimes \log \nabla_p^T[s, u] \leftarrow v \otimes \gamma$ 
             $\gamma \leftarrow \beta[u, t, L, 0] \otimes \alpha[u, s - 1, L, 1] \otimes \theta_{s, t}$ 
             $\nabla_p^T[s, t, L, 0], \nabla_p^T[s, t, L, 1] \leftarrow \gamma \otimes \log \nabla_p^T[s, u] \leftarrow v \otimes \gamma$ 
           $v \leftarrow \nabla_p^T[s, t, R, 1] \otimes \beta[p, t, R, 1]$ 
          for  $u = 1, \dots, s$  do
             $\nabla_p^T[s, u, R, 1], \nabla_p^T[s, u, R, 0] \leftarrow v \otimes \beta[u, u, R, 1] \otimes \alpha[u, s, R, 0]$ 
          for  $u = t + 1, \dots, n$  do
             $v \leftarrow \beta[u, u, R, 0] \otimes \alpha[t + 1, u, L, 1] \otimes \theta_{s, t}$ 
             $\nabla_p^T[s, u, R, 0], \nabla_p^T[s, u, L, 1], \log \nabla_p^T[s, u] \leftarrow v \otimes \gamma$ 
             $\gamma \leftarrow \beta[u, u, L, 0] \otimes \alpha[t + 1, u, L, 1] \otimes \theta_{s, t}$ 
             $\nabla_p^T[s, u, L, 0], \nabla_p^T[s, u, L, 1], \log \nabla_p^T[s, u] \leftarrow v \otimes \gamma$ 
    for  $k = n, \dots, 1$  do
      for  $s = 1, \dots, n - k$  do
         $t \leftarrow s + k$ 
         $v \leftarrow \nabla_p^T[s, t, R, 1] \otimes \alpha[s, t, R, 1]$ 
        for  $u = s + 1, \dots, t$  do
           $\nabla_p^T[s, u, R, 0], \nabla_p^T[s, u, R, 1] \leftarrow v \otimes \alpha[s, u, R, 0] \otimes \alpha[s, t, R, 1]$ 
        if  $s > 1$  then
           $v \leftarrow \nabla_p^T[s, t, L, 1] \otimes \alpha[s, t, L, 1]$ 
          for  $u = s, \dots, t - 1$  do
             $\nabla_p^T[s, u, L, 1], \nabla_p^T[s, u, L, 0] \leftarrow v \otimes \alpha[s, u, L, 1] \otimes \alpha[s, t, L, 0]$ 
           $v \leftarrow \nabla_p^T[s, t, L, 0] \otimes \alpha[s, t, L, 0]$ 
          for  $u = s, \dots, t - 1$  do
             $\gamma \leftarrow \alpha[s, u, R, 1] \otimes \alpha[u + 1, t, L, 1] \otimes \theta_{s, t}$ 
             $\nabla_p^T[s, u, R, 1], \nabla_p^T[s, u, L, 1], \log \nabla_p^T[s, u] \leftarrow v \otimes \gamma$ 
           $v \leftarrow \nabla_p^T[s, t, R, 0] \otimes \alpha[s, t, R, 0]$ 
          for  $u = s, \dots, t - 1$  do
             $\gamma \leftarrow \alpha[s, u, R, 1] \otimes \alpha[u + 1, t, L, 1] \otimes \theta_{s, t}$ 
             $\nabla_p^T[s, u, R, 1], \nabla_p^T[s, u, L, 1], \log \nabla_p^T[s, u] \leftarrow v \otimes \gamma$ 
           $\nabla_p^T[s, u, R, 1], \nabla_p^T[s, u, L, 1], \log \nabla_p^T[s, u] \leftarrow v \otimes \gamma$ 
    return  $\text{sign} \exp \log \nabla_p^T$ 

```

Figure 7: Backpropagation through the inside-outside algorithm to calculate the gradient with respect to the input potentials.  $\nabla_p^T$  denotes the Jacobian of  $\alpha$  with respect to  $\theta$  ( $\nabla_p^T$  is the gradient with respect to  $\theta$ ).  $\alpha, \beta \leftarrow \exp \circ \text{concat} \leftarrow \alpha \otimes \beta$  and  $\beta \leftarrow \beta \otimes \alpha$ .

xact.

inals are dense;  
nation)

ugh DPs is tricky;  
[8])

# Back-propagating through marginals

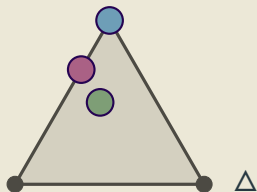
## Pros:

- Familiar algorithms for NLPers,
- (Structured Attention Networks:) All computations exact.

## Cons:


- (Structured Attention Networks:) forward pass marginals are dense;  
(fixed by Perturb & MAP, at cost of rough approximation)
- Efficient & numerically stable back-propagation through DPs is tricky;  
(somewhat alleviated by Mensch and Blondel [2018])
- Not applicable when marginals are unavailable.
- Case-by-case algorithms required, can get tedious.

- **argmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s}$
- **softmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s} + H(\mathbf{p})$
- **sparsemax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s} - 1/2 \|\mathbf{p}\|^2$

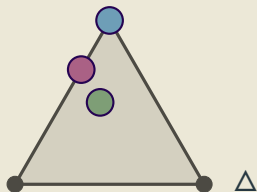


● **MAP**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta}$

● **marginals**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta} + \tilde{H}(\boldsymbol{\mu})$




- **argmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s}$
- **softmax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s} + H(\mathbf{p})$
- **sparsemax**  $\arg \max_{\mathbf{p} \in \Delta} \mathbf{p}^T \mathbf{s} - 1/2 \|\mathbf{p}\|^2$



- **MAP**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta}$

- **marginals**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta} + \tilde{H}(\boldsymbol{\mu})$

- **SparseMAP**  $\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^T \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$



# SparseMAP solution

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

$$= \begin{array}{c} \circ \quad \circ \\ \diagup \quad \diagdown \\ \circ \quad \circ \end{array} = .6 \begin{array}{c} \circ \quad \circ \\ \diagup \quad \diagdown \\ \circ \quad \circ \end{array} + .4 \begin{array}{c} \circ \quad \circ \\ \diagup \quad \diagdown \\ \circ \quad \circ \end{array}$$

( $\boldsymbol{\mu}^*$  is unique, but may have multiple decompositions  $\mathbf{p}$ . Active Set recovers a sparse one.)

# Algorithms for SparseMAP

$$\boldsymbol{\mu}^{\star} = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^{\top} \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

This is also  $\text{proj}_{\mathcal{M}}$  required by SPIGOT!

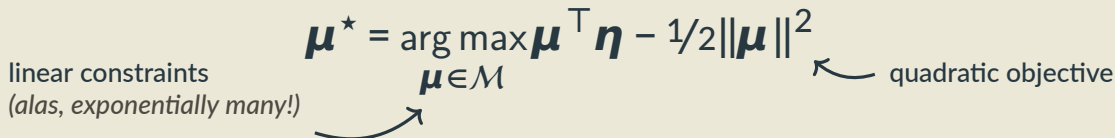


# Algorithms for SparseMAP

linear constraints  
(*alas, exponentially many!*)

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

quadratic objective

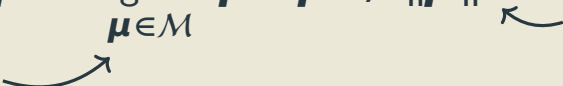


# Algorithms for SparseMAP

linear constraints  
(*alas, exponentially many!*)

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

quadratic objective



## Conditional Gradient

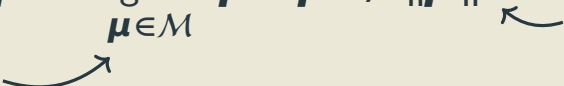
[Frank and Wolfe, 1956, Lacoste-Julien and Jaggi, 2015]

# Algorithms for SparseMAP

linear constraints  
(*alas, exponentially many!*)

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

quadratic objective



## Conditional Gradient

[Frank and Wolfe, 1956, Lacoste-Julien and Jaggi, 2015]

- select a new corner of  $\mathcal{M}$

# Algorithms for SparseMAP

linear constraints  
(*alas, exponentially many!*)

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

quadratic objective

## Conditional Gradient

[Frank and Wolfe, 1956, Lacoste-Julien and Jaggi, 2015]

- select a new corner of  $\mathcal{M}$

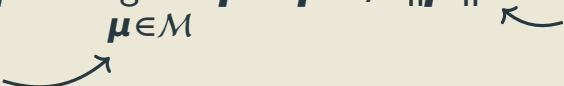
$$\arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \underbrace{(\boldsymbol{\eta} - \boldsymbol{\mu}^{(t-1)})}_{\tilde{\boldsymbol{\eta}}}$$

# Algorithms for SparseMAP

linear constraints  
(*alas, exponentially many!*)

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

quadratic objective



## Conditional Gradient

[Frank and Wolfe, 1956, Lacoste-Julien and Jaggi, 2015]

- select a new corner of  $\mathcal{M}$
- update the (sparse) coefficients of  $\boldsymbol{p}$ 
  - Update rules: vanilla, away-step, pairwise

# Algorithms for SparseMAP

linear constraints  
(*alas, exponentially many!*)

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

quadratic objective

## Conditional Gradient

[Frank and Wolfe, 1956, Lacoste-Julien and Jaggi, 2015]

- select a new corner of  $\mathcal{M}$
- update the (sparse) coefficients of  $\boldsymbol{p}$ 
  - Update rules: vanilla, away-step, pairwise
  - Quadratic objective: **Active Set**

a.k.a. Min-Norm Point, [Wolfe, 1976]

[Martins et al., 2015, Nocedal and Wright, 1999]

# Algorithms for SparseMAP

linear constraints  
(*alas, exponentially many!*)

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

quadratic objective

## Conditional Gradient

[Frank and Wolfe, 1956, Lacoste-Julien and Jaggi, 2015]

- select a new corner
- update the (sparse)
  - Update rules: vanilla
  - Quadratic objective:

Active Set achieves  
**finite & linear** convergence!

a.k.a. Min-Norm Point, [Wolfe, 1976]

[Martins et al., 2015, Nocedal and Wright, 1999]

# Algorithms for SparseMAP

linear constraints  
(*alas, exponentially many!*)

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

quadratic objective

## Conditional Gradient

[Frank and Wolfe, 1956, Lacoste-Julien and Jaggi, 2015]

- select a new corner of  $\mathcal{M}$
- update the (sparse) coefficients of  $\boldsymbol{p}$ 
  - Update rules: vanilla, away-step, pairwise
  - Quadratic objective: **Active Set**  
a.k.a. Min-Norm Point, [Wolfe, 1976]  
[Martins et al., 2015, Nocedal and Wright, 1999]

## Backward pass

$$\frac{\partial \boldsymbol{\mu}}{\partial \boldsymbol{\eta}} \text{ is sparse}$$



# Algorithms for SparseMAP

linear constraints  
(*alas, exponentially many!*)

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

quadratic objective

## Conditional Gradient

[Frank and Wolfe, 1956, Lacoste-Julien and Jaggi, 2015]

- select a new corner of  $\mathcal{M}$
- update the (sparse) coefficients of  $\mathbf{p}$ 
  - Update rules: vanilla, away-step, pairwise
  - Quadratic objective: **Active Set**  
a.k.a. Min-Norm Point, [Wolfe, 1976]

[Martins et al., 2015, Nocedal and Wright, 1999]

## Backward pass

$\frac{\partial \boldsymbol{\mu}}{\partial \boldsymbol{\eta}}$  is sparse  
computing  $\left(\frac{\partial \boldsymbol{\mu}}{\partial \boldsymbol{\eta}}\right)^\top d\mathbf{y}$   
takes  $\mathcal{O}(\dim(\boldsymbol{\mu}) \text{nnz}(\mathbf{p}^*))$

# Algorithms for SparseMAP

$$\boldsymbol{\mu}^* = \arg \max_{\boldsymbol{\mu} \in \mathcal{M}} \boldsymbol{\mu}^\top \boldsymbol{\eta} - 1/2 \|\boldsymbol{\mu}\|^2$$

linear constraints  
(*alas, exponentially many!*)

quadratic objective

Condition

pass

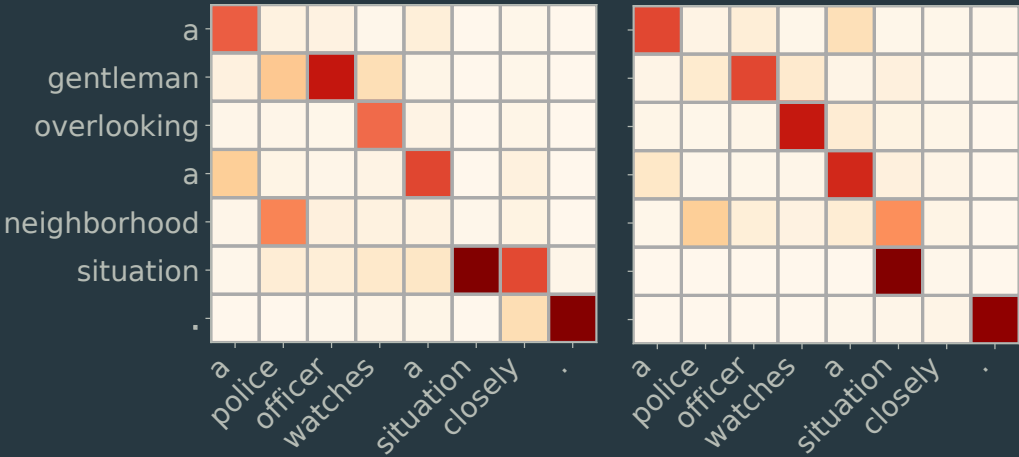
Completely modular: just add MAP



[Frank and Wolfe, 1956]

- select a new c
- update the (sparse) coefficients of  $\mathbf{p}$ 
  - Update rules: vanilla, away-step, pairwise
  - Quadratic objective: **Active Set**  
a.k.a. Min-Norm Point, [Wolfe, 1976]

[Martins et al., 2015, Nocedal and Wright, 1999]

computing  $\left(\frac{\partial \boldsymbol{\mu}}{\partial \boldsymbol{\eta}}\right)^\top d\boldsymbol{\eta}$   
takes  $\mathcal{O}(\dim(\boldsymbol{\mu}) \text{nnz}(\mathbf{p}^*))$





# Overview

$$\mathbb{E}_{\pi_{\boldsymbol{\theta}}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- REINFORCE
- Straight-Through Gumbel (Perturb & MAP)

$$L\left(\arg \max_{\mathbf{z}} \pi_{\boldsymbol{\theta}}(\mathbf{z} | x)\right)$$

- Straight-Through
- SPIGOT

$$L\left(\mathbb{E}_{\pi_{\boldsymbol{\theta}}(\mathbf{z}|x)}[\mathbf{z}]\right)$$

- Structured Attn. Nets
- SparseMAP

# Structured latent variables without sampling

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}(\mathbf{z})) \pi(\mathbf{z} | \mathbf{x})$$

# Structured latent variables without sampling

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\boldsymbol{\phi}}(\mathbf{z})) \pi_{\boldsymbol{\theta}}(\mathbf{z} \mid \mathbf{x})$$

# Structured latent variables without sampling

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | \mathbf{x})$$

e.g., a TreeLSTM defined by  $\mathbf{z}$



# Structured latent variables without sampling

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\boldsymbol{\phi}}(\mathbf{z})) \pi_{\boldsymbol{\theta}}(\mathbf{z} | \mathbf{x})$$

e.g., a TreeLSTM defined by  $\mathbf{z}$

parsing model,  
using some scorer  $f_{\boldsymbol{\theta}}(\mathbf{z}; \mathbf{x})$

# Structured latent variables without sampling

sum over  
all possible trees

e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\boldsymbol{\phi}}(\mathbf{z})) \pi_{\boldsymbol{\theta}}(\mathbf{z} | \mathbf{x})$$

parsing model,  
using some scorer  $\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{z}; \mathbf{x})$

Exponentially large sum!

# Structured latent variables without sampling

sum over  
all possible trees

e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | \mathbf{x})$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; \mathbf{x})$

idea 1

idea 2

idea 3

# Structured latent variables without sampling

sum over  
all possible trees

e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | x)$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; x)$

$$\sum_{h \in \mathcal{H}}$$

idea 1

idea 2

idea 3

# Structured latent variables without sampling

sum over  
all possible trees

e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | x)$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; x)$

$$\sum_{h \in \mathcal{H}} \frac{\partial \mathbb{E}[L(\mathbf{z})]}{\partial \theta}$$

idea 1

idea 2

idea 3

# Structured latent variables without sampling

sum over  
all possible trees

e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | x)$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; x)$

$$\sum_{h \in \mathcal{H}} \frac{\partial \mathbb{E}[L(\mathbf{z})]}{\partial \theta}$$

idea 1  $\pi_{\theta}(\mathbf{z}) \propto \exp(f_{\theta}(\mathbf{z}))$

softmax

idea 2

idea 3

# Structured latent variables without sampling

sum over  
all possible trees


e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | x)$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; x)$

$$\sum_{h \in \mathcal{H}} \frac{\partial \mathbb{E}[L(\mathbf{z})]}{\partial \theta}$$

softmax 

idea 1  $\pi_{\theta}(\mathbf{z}) \propto \exp(f_{\theta}(\mathbf{z}))$

idea 2

idea 3

# Structured latent variables without sampling

sum over  
all possible trees

e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | x)$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; x)$

$$\sum_{h \in \mathcal{H}} \frac{\partial \mathbb{E}[L(\mathbf{z})]}{\partial \theta}$$

softmax 🙌 😊

idea 1  $\pi_{\theta}(\mathbf{z}) \propto \exp(f_{\theta}(\mathbf{z}))$

idea 2

idea 3



# Structured latent variables without sampling

sum over  
all possible trees

e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | \mathbf{x})$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; \mathbf{x})$

All methods we've seen require sampling; hard in general.

idea 2

idea 3

# Structured latent variables without sampling

sum over  
all possible trees



e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | x)$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; x)$

$$\sum_{h \in \mathcal{H}} \frac{\partial \mathbb{E}[L(\mathbf{z})]}{\partial \theta}$$

idea 1  $\pi_{\theta}(\mathbf{z}) \propto \exp(f_{\theta}(\mathbf{z}))$

softmax

idea 2  $\pi_{\theta}(\mathbf{z}) = 1$  if  $\mathbf{z} = \text{MAP}(f_{\theta}(\cdot))$  else 0

argmax

idea 3

# Structured latent variables without sampling

sum over  
all possible trees

e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | x)$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; x)$

$$\sum_{h \in \mathcal{H}} \frac{\partial \mathbb{E}[L(\mathbf{z})]}{\partial \theta}$$

idea 1  $\pi_{\theta}(\mathbf{z}) \propto \exp(f_{\theta}(\mathbf{z}))$

softmax



idea 2  $\pi_{\theta}(\mathbf{z}) = 1$  if  $\mathbf{z} = \text{MAP}(f_{\theta}(\cdot))$  else 0

argmax



idea 3

# Structured latent variables without sampling

sum over  
all possible trees

e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | x)$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; x)$

$$\sum_{h \in \mathcal{H}} \frac{\partial \mathbb{E}[L(\mathbf{z})]}{\partial \theta}$$

idea 1  $\pi_{\theta}(\mathbf{z}) \propto \exp(f_{\theta}(\mathbf{z}))$

softmax



idea 2  $\pi_{\theta}(\mathbf{z}) = 1$  if  $\mathbf{z} = \text{MAP}(f_{\theta}(\cdot))$  else 0

argmax



idea 3

# Structured latent variables without sampling

sum over  
all possible trees

e.g., a TreeLSTM defined by  $\mathbf{z}$

$$\mathbb{E}_{\mathbf{z}}[L(\mathbf{z})] = \sum_{\mathbf{z} \in \mathcal{Z}} L(\hat{\mathbf{y}}_{\phi}(\mathbf{z})) \pi_{\theta}(\mathbf{z} | x)$$

How to define  $\pi_{\theta}$ ?

parsing model,  
using some scorer  $f_{\theta}(\mathbf{z}; x)$

$$\sum_{h \in \mathcal{H}} \frac{\partial \mathbb{E}[L(\mathbf{z})]}{\partial \theta}$$

idea 1  $\pi_{\theta}(\mathbf{z}) \propto \exp(f_{\theta}(\mathbf{z}))$

idea 2  $\pi_{\theta}(\mathbf{z}) = 1$  if  $\mathbf{z} = \text{MAP}(f_{\theta}(\cdot))$  else 0

idea 3

softmax

argmax

SparseMAP



# Structured latent variables without sampling

$$\begin{array}{c} \text{•} \quad \text{•} \quad \text{•} \\ \text{↖} \quad \text{↗} \\ \text{•} \end{array} = .7x \quad \begin{array}{c} \text{•} \quad \text{•} \quad \text{•} \\ \text{↖} \quad \text{↗} \\ \text{•} \end{array} + .3x \quad \begin{array}{c} \text{•} \quad \text{•} \quad \text{•} \\ \text{↖} \quad \text{↗} \\ \text{•} \end{array}$$

# Structured latent variables without sampling

$$\begin{array}{c} \text{•} \quad \text{•} \quad \text{•} \\ \text{↖} \quad \text{↗} \\ \text{•} \end{array} = .7x \quad \begin{array}{c} \text{•} \quad \text{•} \quad \text{•} \\ \text{↖} \quad \text{↗} \\ \text{•} \end{array} + .3x \quad \begin{array}{c} \text{•} \quad \text{•} \quad \text{•} \\ \text{↖} \quad \text{↗} \\ \text{•} \end{array} + 0x \quad \begin{array}{c} \text{•} \quad \text{•} \quad \text{•} \\ \text{↖} \quad \text{↗} \\ \text{•} \end{array} + \dots$$

# Structured latent variables without sampling

$$\begin{aligned} \text{Diagram} &= .7 \times \text{Diagram} + .3 \times \text{Diagram} + 0 \times \text{Diagram} + \dots \\ \mathbb{E}[L(\mathbf{z})] &= .7 \times L(\text{Diagram}) + .3 \times L(\text{Diagram}) \end{aligned}$$

recall our shorthand  $L(\mathbf{z}) = L(\hat{y}_{\phi}(\mathbf{z}), y)$



# V. Conclusions

## Stanford Natural Language Inference (Accuracy)

[Kim et al., 2017]

Simple Attention 86.2

Structured Attention 86.8

[Liu and Lapata, 2018]

100D Structured Attention 86.8

[Yogatama et al., 2017]

100D RL-SPINN 80.5

[Choi et al., 2018]

100D ST Gumbel-Tree 82.6

300D - 85.6

600D - 86.0

[Corro and Titov, 2019b]

Latent Tree + 1 GCN - 85.2

Latent Tree + 2 GCN - 86.2

[Havrylov et al., 2019]

100D TreeLSTM + tricks 84.3

## Stanford Sentiment (Accuracy)

[Socher et al., 2013]

Bigram Naive Bayes 83.1

[Niculae et al., 2018b]

DepTreeLSTM w/ CoreNLP 83.2

DepTreeLSTM w/ SparseMAP 84.7

[Corro and Titov, 2019b]

GCN w/ CoreNLP 83.8

GCN w/ Perturb-and-MAP 84.6

[Choi et al., 2018]

ST Gumbel-Tree 90.7

[Havrylov et al., 2019]

TreeLSTM + tricks 90.2

# Is it syntax?!


- Unlike e.g. unsupervised parsing, the structures we learn are guided by a **downstream objective** (typically discriminative).
- They don't typically resemble grammatical structure (yet) [Williams et al., 2018]  
(future work: more inductive biases and constraints?)

# Is it syntax?!

- Unlike e.g. unsupervised parsing, the structures we learn are guided by a **downstream objective** (typically discriminative).
- They don't typically resemble grammatical structure (yet) [Williams et al., 2018]  
(future work: more inductive biases and constraints?)
- Common to compare latent structures with parser outputs.  
But is this always a meaningful comparison?

# Syntax vs. Composition Order

CoreNLP parse,  $p = 21.4\%$




# Syntax vs. Composition Order

$p = 22.6\%$

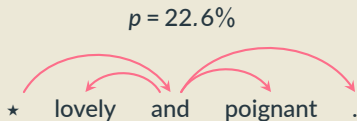


CoreNLP parse,  $p = 21.4\%$


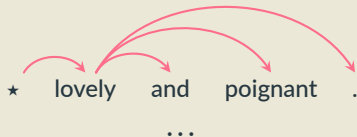


...

# Syntax vs. Composition Order




CoreNLP parse,  $p = 21.4\%$



$p = 15.27\%$



...  
CoreNLP parse,  $p = 0\%$



# Conclusions

- Latent structure models are desirable for interpretability, structural bias, and higher predictive power with fewer parameters.
- Stochastic latent variables can be dealt with RL or straight-through gradients.
- Deterministic argmax requires surrogate gradients (e.g. SPIGOT).
- Continuous relaxations of argmax include SANs and SparseMAP.
- Intuitively, some of these different methods are trying to do similar things or require the same building blocks (e.g. SPIGOT and SparseMAP).
- ... we didn't even get into deep *generative* models! These tools apply, but there are new challenges. [Corro and Titov, 2019a, Kim et al., 2019a,b, Kawakami et al., 2019]



# Overview

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

- REINFORCE
- Straight-Through Gumbel (Perturb & MAP)
- SparseMAP

$$L(\arg \max_{\mathbf{z}} \pi_{\theta}(\mathbf{z} | x))$$

- Straight-Through
- SPIGOT

$$L(\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[\mathbf{z}])$$

- Structured Attn. Nets
- SparseMAP

# Overview

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

$$L\left(\arg \max_{\mathbf{z}} \pi_{\theta}(\mathbf{z} | x)\right)$$

$$L\left(\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[\mathbf{z}]\right)$$

- REINFORCE<sup>SPL</sup>
- Straight-Through Gumbel  
(Perturb & MAP)<sup>SPL,MRG</sup>
- SparseMAP<sup>MAP+</sup>

- Straight-Through<sup>MAP,MRG</sup>
- SPIGOT<sup>MAP+</sup>

- Structured Attn. Nets<sup>MRG</sup>
- SparseMAP<sup>MAP+</sup>

## Computation:

<sup>SPL</sup>: Sampling. (Simple in incremental/unstructured, hard for most global structures.)

<sup>MAP</sup>: Finding the highest-scoring structure.

<sup>MRG</sup>: Marginal inference.

# Overview

Thank you!

$$\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[L(\mathbf{z})]$$

$$L(\arg \max_{\mathbf{z}} \pi_{\theta}(\mathbf{z} | x))$$

$$L(\mathbb{E}_{\pi_{\theta}(\mathbf{z}|x)}[\mathbf{z}])$$

- REINFORCE<sup>SPL</sup>
- Straight-Through Gumbel  
(Perturb & MAP)<sup>SPL,MRG</sup>
- SparseMAP<sup>MAP+</sup>

- Straight-Through<sup>MAP,MRG</sup>
- SPIGOT<sup>MAP+</sup>

- Structured Attn. Nets<sup>MRG</sup>
- SparseMAP<sup>MAP+</sup>

## Computation:

<sup>SPL</sup>: Sampling. (Simple in incremental/unstructured, hard for most global structures.)

<sup>MAP</sup>: Finding the highest-scoring structure.

<sup>MRG</sup>: Marginal inference.

# References I

- Ryan Adams. The gumbel-max trick for discrete distributions, 2013. URL <https://lips.cs.princeton.edu/the-gumbel-max-trick-for-discrete-distributions/>. Blog post.
- James K Baker. Trainable grammars for speech recognition. *The Journal of the Acoustical Society of America*, 65(S1):S132–S132, 1979.
- Yoshua Bengio, Nicholas Léonard, and Aaron Courville. Estimating or propagating gradients through stochastic neurons for conditional computation. *arXiv preprint arXiv:1308.3432*, 2013.
- Mathieu Blondel, André FT Martins, and Vlad Niculae. Learning classifiers with Fenchel-Young losses: Generalized entropies, margins, and algorithms. In *Proc. of AISTATS*, 2019.
- Samuel R. Bowman, Jon Gauthier, Abhinav Rastogi, Raghav Gupta, Christopher D. Manning, and Christopher Potts. A fast unified model for parsing and sentence understanding. In *Proc. of ACL*, 2016. doi: 10.18653/v1/P16-1139.
- Stephen Boyd and Lieven Vandenberghe. *Convex optimization*. Cambridge University Press, 2004.
- Peter F Brown, Vincent J Della Pietra, Stephen A Della Pietra, and Robert L Mercer. The mathematics of statistical machine translation: Parameter estimation. *Computational Linguistics*, 19(2):263–311, 1993.
- Qian Chen, Xiaodan Zhu, Zhen-Hua Ling, Si Wei, Hui Jiang, and Diana Inkpen. Enhanced LSTM for natural language inference. In *Proc. of ACL*, 2017.
- Jihun Choi, Kang Min Yoo, and Sang-goo Lee. Learning to compose task-specific tree structures. In *Proc. of AAAI*, 2018.

# References II

- Yoeng-Jin Chu and Tseng-Hong Liu. On the shortest arborescence of a directed graph. *Science Sinica*, 14:1396–1400, 1965.
- William John Cocke and Jacob T Schwartz. *Programming languages and their compilers*. Courant Institute of Mathematical Sciences., 1970.
- Shay B Cohen, Karl Stratos, Michael Collins, Dean P Foster, and Lyle Ungar. Spectral learning of latent-variable PCFGs. In *Proc. of ACL*, 2012.
- Calo M Correia, Gon Vlad Niculae, and André FT Martins. Adaptively sparse transformers. In *Proc. of EMNLP-IJCNLP (to appear)*, 2019.
- Caio Corro and Ivan Titov. Differentiable Perturb-and-Parse: Semi-Supervised Parsing with a Structured Variational Autoencoder. In *Proc. of ICLR*, 2019a.
- Caio Corro and Ivan Titov. Learning latent trees with stochastic perturbations and differentiable dynamic programming. In *Proc. of ACL*, 2019b.
- Marco Cuturi and Mathieu Blondel. Soft-DTW: a differentiable loss function for time-series. In *Proc. of ICML*, 2017.
- Jack Edmonds. Optimum branchings. *J. Res. Nat. Bur. Stand.*, 71B:233–240, 1967.
- Marguerite Frank and Philip Wolfe. An algorithm for quadratic programming. *Nav. Res. Log.*, 3(1-2):95–110, 1956.
- Serhii Havrylov, Germán Kruszewski, and Armand Joulin. Cooperative Learning of Disjoint Syntax and Semantics. In *Proc. NAACL-HLT*, 2019.

# References III

- Geoffrey Hinton. Neural networks for machine learning. In *Coursera video lectures*, 2012.
- Eric Jang, Shixiang Gu, and Ben Poole. Categorical reparameterization with Gumbel-softmax. In *Proc. of ICLR*, 2017.
- Roy Jonker and Anton Volgenant. A shortest augmenting path algorithm for dense and sparse linear assignment problems. *Computing*, 38(4):325–340, 1987.
- Tadao Kasami. An efficient recognition and syntax-analysis algorithm for context-free languages. *Coordinated Science Laboratory Report no. R-257*, 1966.
- Kazuya Kawakami, Chris Dyer, and Phil Blunsom. Learning to discover, ground and use words with segmental neural language models. In *Proc. of ACL*, 2019.
- Yoon Kim, Carl Denton, Loung Hoang, and Alexander M Rush. Structured attention networks. In *Proc. of ICLR*, 2017.
- Yoon Kim, Chris Dyer, and Alexander Rush. Compound probabilistic context-free grammars for grammar induction. In *Proc. of ACL*, 2019a.
- Yoon Kim, Alexander Rush, Lei Yu, Adhiguna Kuncoro, Chris Dyer, and Gábor Melis. Unsupervised recurrent neural network grammars. In *Proc. of NAACL-HLT*, 2019b.
- Diederik P Kingma and Max Welling. Auto-encoding Variational Bayes. 2014.
- Gustav Kirchhoff. Ueber die auflösung der gleichungen, auf welche man bei der untersuchung der linearen vertheilung galvanischer ströme geführt wird. *Annalen der Physik*, 148(12):497–508, 1847.

# References IV

- Harold W Kuhn. The Hungarian method for the assignment problem. *Nav. Res. Log.*, 2(1-2):83–97, 1955.
- Simon Lacoste-Julien and Martin Jaggi. On the global linear convergence of Frank-Wolfe optimization variants. In *Proc. of NeurIPS*, 2015.
- Zhifei Li and Jason Eisner. First-and second-order expectation semirings with applications to minimum-risk training on translation forests. In *Proc. of EMNLP*, 2009.
- Yang Liu and Mirella Lapata. Learning structured text representations. *TACL*, 6:63–75, 2018.
- Chris J Maddison, Andriy Mnih, and Yee Whye Teh. The concrete distribution: A continuous relaxation of discrete random variables. In *Proc. of ICLR*, 2016.
- Jean Maillard and Stephen Clark. Latent tree learning with differentiable parsers: Shift-Reduce parsing and chart parsing. *arXiv preprint arXiv:1806.00840*, 2018.
- Chaitanya Malaviya, Pedro Ferreira, and André FT Martins. Sparse and constrained attention for neural machine translation. In *Proc. of ACL*, 2018.
- André FT Martins and Ramón Fernandez Astudillo. From softmax to sparsemax: A sparse model of attention and multi-label classification. In *Proc. of ICML*, 2016.
- André FT Martins and Julia Kreutzer. Learning what’s easy: Fully differentiable neural easy-first taggers. In *Proc. of EMNLP*, 2017.
- André FT Martins and Vlad Niculae. Notes on latent structure models and SPIGOT. *preprint arXiv:1907.10348*, 2019.

# References V

- André FT Martins, Mário AT Figueiredo, Pedro MQ Aguiar, Noah A Smith, and Eric P Xing. AD3: Alternating directions dual decomposition for MAP inference in graphical models. *JMLR*, 16(1):495–545, 2015.
- Arthur Mensch and Mathieu Blondel. Differentiable dynamic programming for structured prediction and attention. In *Proc. of ICML*, 2018.
- Nikita Nangia and Samuel Bowman. ListOps: A diagnostic dataset for latent tree learning. In *Proc. of NAACL SRW*, 2018.
- Vlad Niculae and Mathieu Blondel. A regularized framework for sparse and structured neural attention. In *Proc. of NeurIPS*, 2017.
- Vlad Niculae, André FT Martins, Mathieu Blondel, and Claire Cardie. SparseMAP: Differentiable sparse structured inference. In *Proc. of ICML*, 2018a.
- Vlad Niculae, André FT Martins, and Claire Cardie. Towards dynamic computation graphs via sparse latent structure. In *Proc. of EMNLP*, 2018b.
- Jorge Nocedal and Stephen Wright. *Numerical Optimization*. Springer New York, 1999.
- George Papandreou and Alan L Yuille. Perturb-and-MAP random fields: Using discrete optimization to learn and sample from energy models. In *Proc. of ICCV*, 2011.
- Hao Peng, Sam Thomson, and Noah A Smith. Backpropagating through structured argmax using a SPIGOT. In *Proc. of ACL*, 2018.
- Ben Peters, Vlad Niculae, and André FT Martins. Sparse sequence-to-sequence models. In *Proc. of ACL*, 2019.



# References VI

- Slav Petrov and Dan Klein. Discriminative log-linear grammars with latent variables. In *Advances in neural information processing systems*, pages 1153–1160, 2008.
- Ariadna Quattoni, Sybor Wang, Louis-Philippe Morency, Michael Collins, and Trevor Darrell. Hidden conditional random fields. *IEEE Transactions on Pattern Analysis & Machine Intelligence*, 29(10):1848–1852, 2007.
- Lawrence R. Rabiner. A tutorial on Hidden Markov Models and selected applications in speech recognition. *P. IEEE*, 77(2):257–286, 1989.
- Hiroaki Sakoe and Seibi Chiba. Dynamic programming algorithm optimization for spoken word recognition. *IEEE Trans. on Acoustics, Speech, and Sig. Proc.*, 26:43–49, 1978.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D Manning, Andrew Ng, and Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proc. EMNLP*, 2013.
- Veselin Stoyanov, Alexander Ropson, and Jason Eisner. Empirical risk minimization of graphical model parameters given approximate inference, decoding, and model structure. In *Proc. of AISTATS*, 2011.
- Ben Taskar. *Learning structured prediction models: A large margin approach*. PhD thesis, Stanford University, 2004.
- Constantino Tsallis. Possible generalization of Boltzmann-Gibbs statistics. *Journal of Statistical Physics*, 52:479–487, 1988.
- Leslie G Valiant. The complexity of computing the permanent. *Theor. Comput. Sci.*, 8(2):189–201, 1979.

# References VII

Tim Vieira. Gumbel-max trick, 2014. URL

<https://timvieira.github.io/blog/post/2014/07/31/gumbel-max-trick/>. Blog post.

Martin J Wainwright and Michael I Jordan. *Graphical models, exponential families, and variational inference.*, volume 1. Now Publishers, Inc., 2008.

Adina Williams, Andrew Drozdov, and Samuel R Bowman. Do latent tree learning models identify meaningful structure in sentences? *TACL*, 2018.

Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Mach. Learn.*, 8, 1992.

Philip Wolfe. Finding the nearest point in a polytope. *Mathematical Programming*, 11(1):128–149, 1976.

Dani Yogatama, Phil Blunsom, Chris Dyer, Edward Grefenstette, and Wang Ling. Learning to compose words into sentences with reinforcement learning. In *Proc. of ICLR*, 2017.

Daniel H Younger. Recognition and parsing of context-free languages in time  $n^3$ . *Information and Control*, 10(2):189–208, 1967.